

академия  
больших  
данных



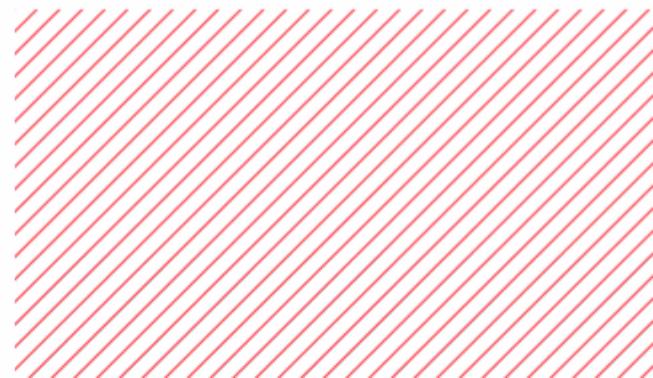
# Высокопроизводительные вычисления. Лекция 3. OpenMP

Рыкованов Сергей

доцент Сколковского института науки и технологий

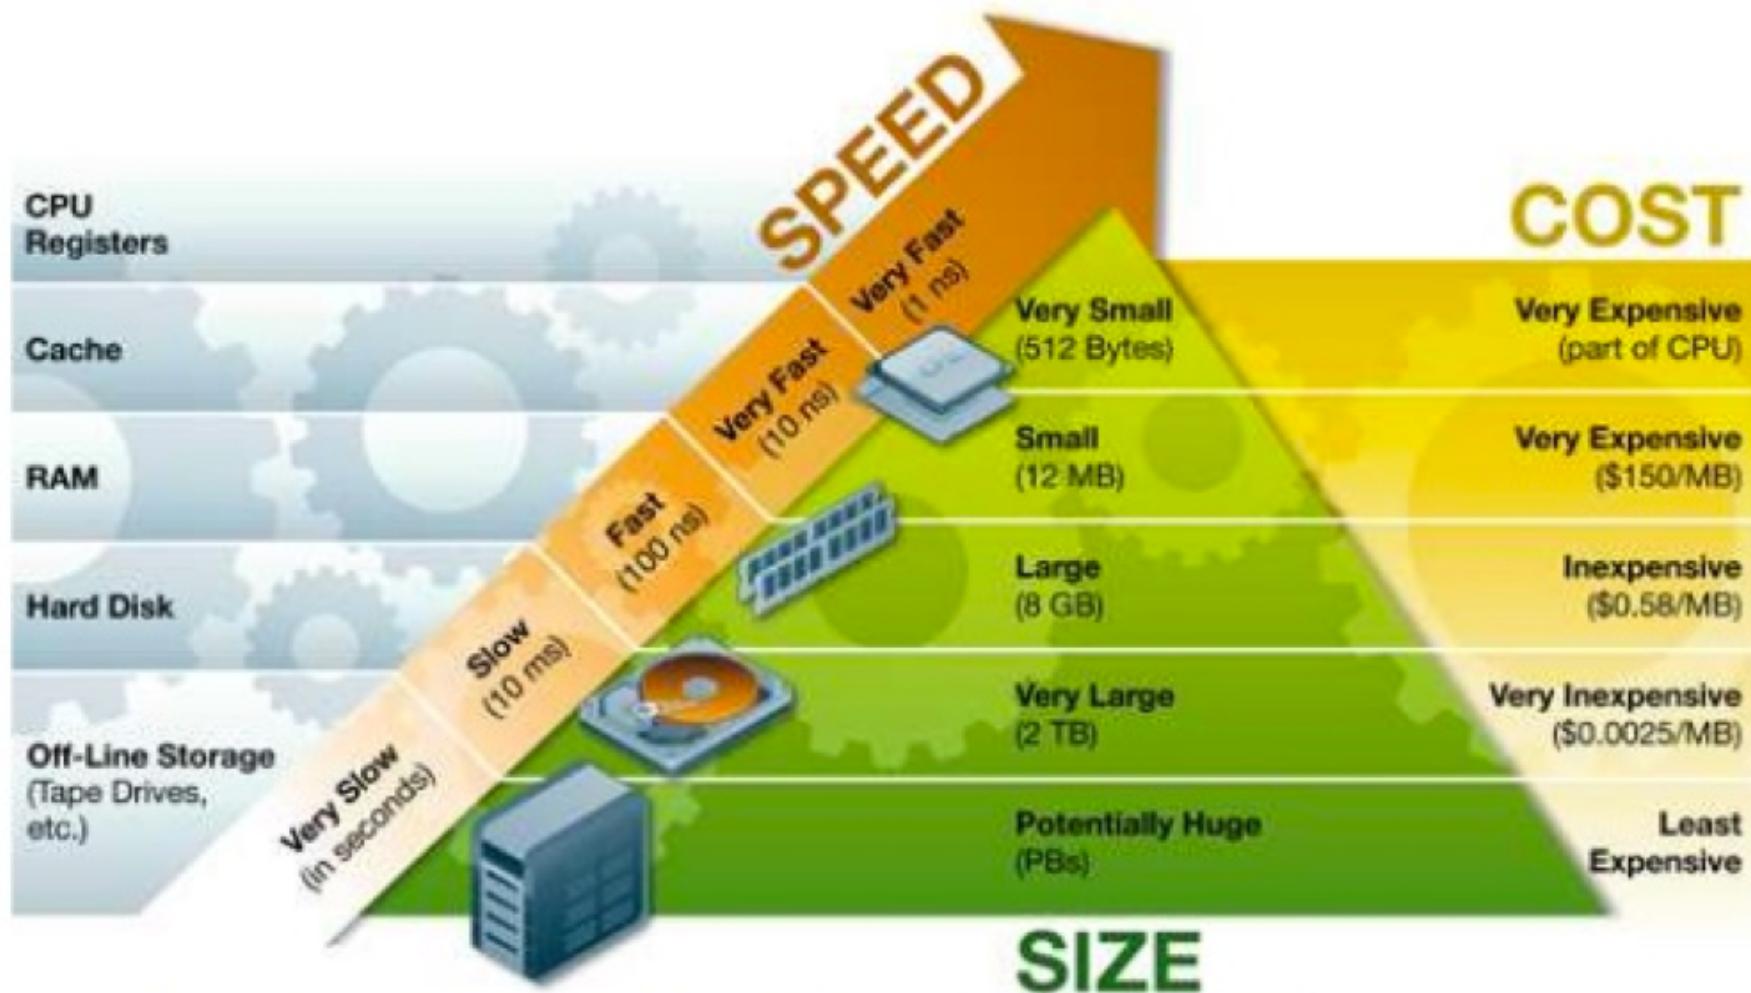
Матвеев Сергей

научный сотрудник Сколковского института науки и технологий



# 1. Перемножение матриц, иерархия памяти

# Иерархия памяти в компьютере



Source: [http://www.ts.avnet.com/uk/products\\_and\\_solutions/storage/hierarchy.html](http://www.ts.avnet.com/uk/products_and_solutions/storage/hierarchy.html)

# «Найвное» перемножение матриц

$$\begin{matrix} c_{00} & c_{01} & c_{02} \\ c_{10} & c_{11} & c_{12} \\ c_{20} & c_{21} & c_{22} \end{matrix} = \begin{matrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \end{matrix} \times \begin{matrix} b_{00} & b_{01} & b_{02} \\ b_{10} & b_{11} & b_{12} \\ b_{20} & b_{21} & b_{22} \end{matrix}$$

NxN matrices: A, B, C

$N \gg 1$

$C = A * B$

$$C_{ij} = \sum_{n=0}^{N-1} a_{in} * b_{nj}$$

```
for i from 0 to N-1:  
    for j from 0 to N-1:  
        for n from 0 to N-1:  
  
            C[i][j] += A[i][n] * B[n][j]
```

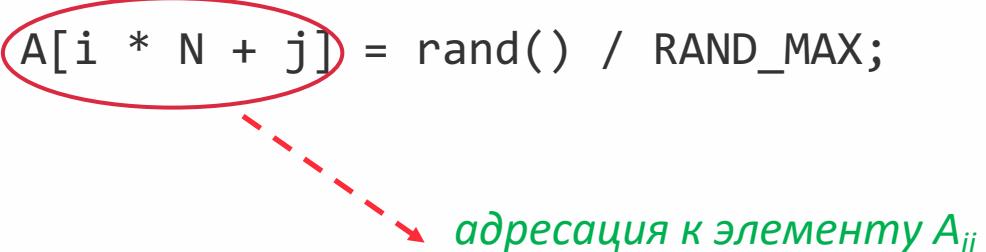
# Основные элементы кода: аллокация памяти и адресация

```
double *A, *B, *C;  
A = (double *) malloc(N * N * sizeof(double));  
B = (double *) malloc(N * N * sizeof(double));  
C = (double *) malloc(N * N * sizeof(double));  
free(A); free(B); free(C);
```



матрица  $N \times N$

```
void RandomMatrix(double * A, size_t N)  
{  
    srand(time(NULL));  
    for (int i = 0; i < N; i++)  
    {  
        for (int j = 0; j < N; j++)  
        {  
            A[i * N + j] = rand() / RAND_MAX;  
        }  
    }  
}
```



адресация к элементу  $A_{ij}$

# Основные элементы кода: matmul и измерение времени

*matmul:*

```
for (i = 0; i < N; i++)
    for(j = 0; j < N; j++)
        for(k = 0; k < N; k++)
            C[i * N + j] = C[i * N + j] + A[i * N + k] * B[k * N + j];
```

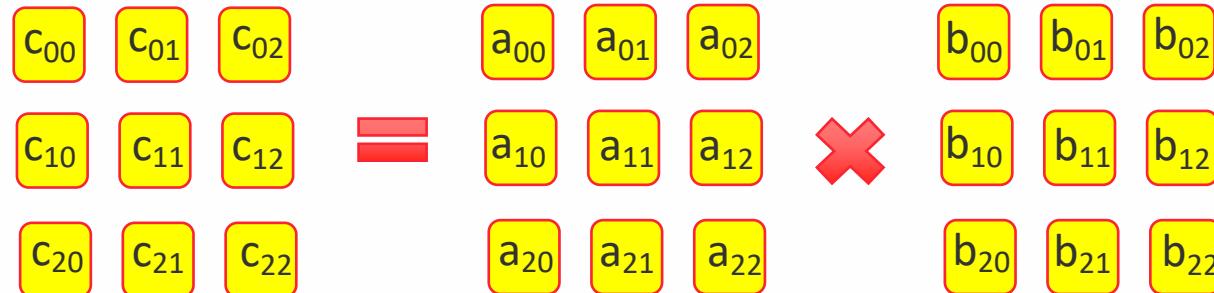
```
struct timeval start, end;
double r_time = 0.0;
gettimeofday(&start, NULL);
... работа
... еще работа
gettimeofday(&end, NULL);
r_time = end.tv_sec - start.tv_sec + ((double)(end.tv_usec - start.tv_usec)) / 1000000;
```



секунды

микросекунды

другие опции с разным разрешением: `clock()`; `clock_gettime()`



```

for i from 0 to N-1:
    for j from 0 to N-1:
        for n from 0 to N-1:

            C[i][j] += A[i][n] * B[n][j]

```

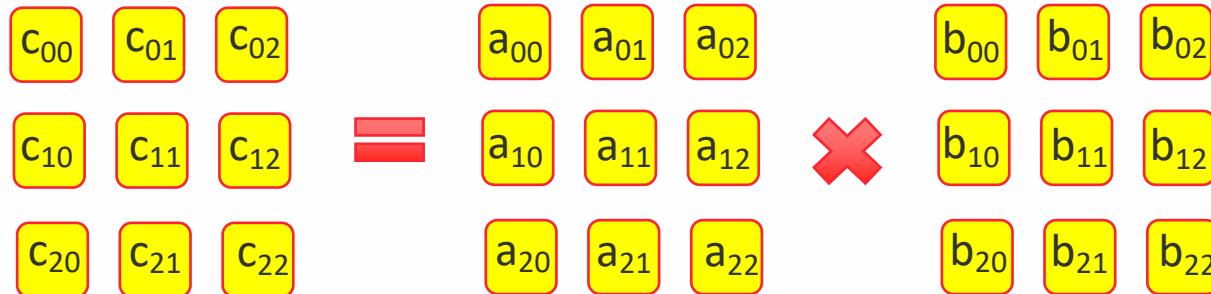
Registers



Cache L1



**RAM**

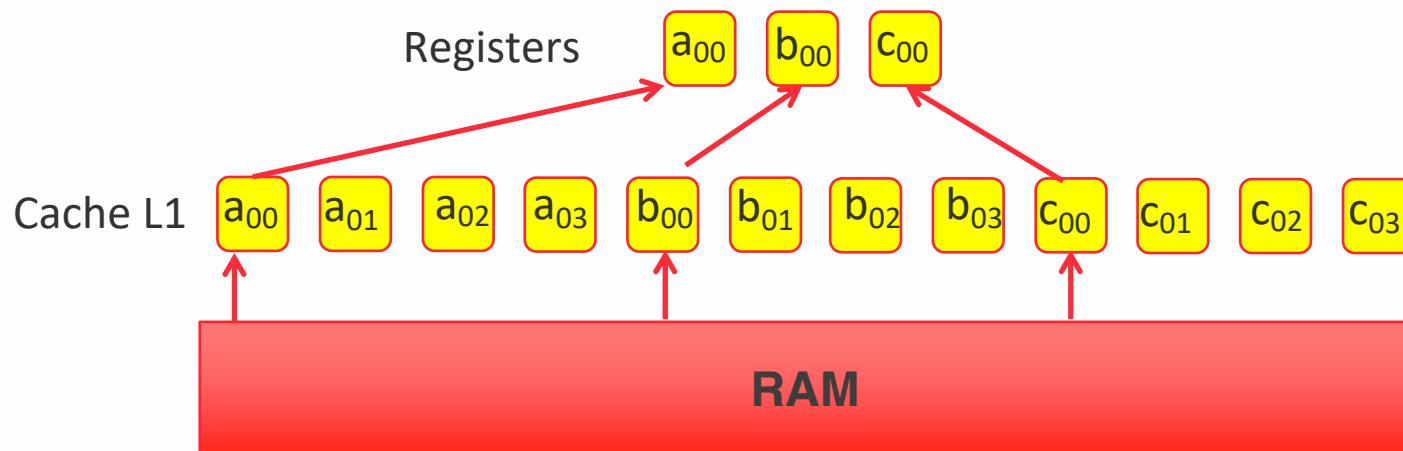


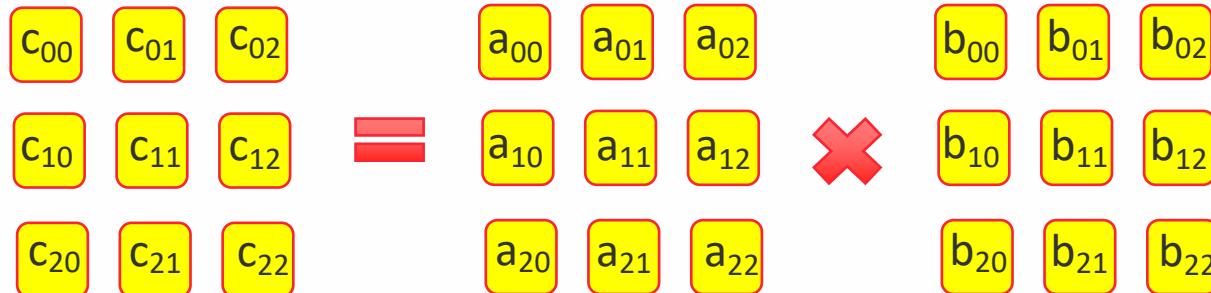
```

for i from 0 to N-1:
    for j from 0 to N-1:
        for n from 0 to N-1:
            C[i][j] += A[i][n]*B[n][j]
    
```

```

i=0
j=0
n=0
C[0][0] += A[0][0]*B[0][0]
    
```



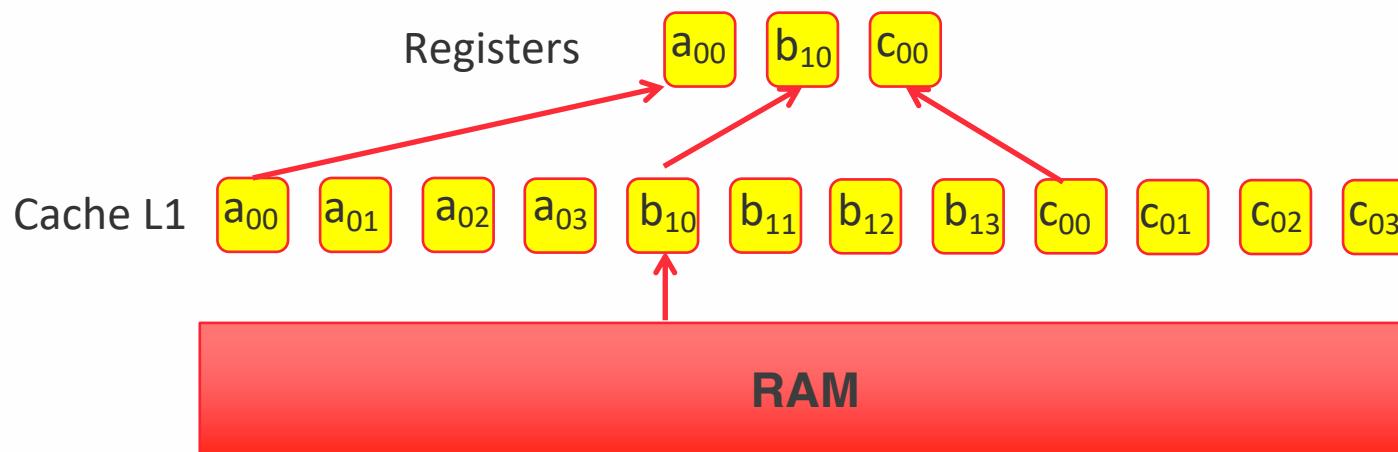


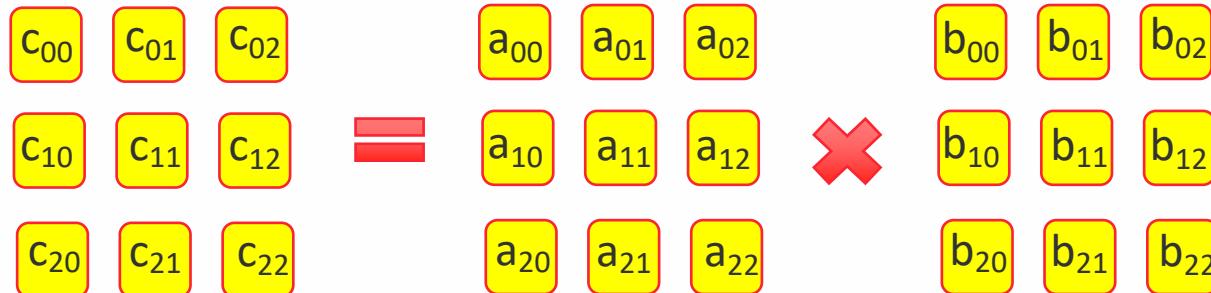
```

for i from 0 to N-1:
    for j from 0 to N-1:
        for n from 0 to N-1:
            C[i][j] += A[i][n]*B[n][j]
    
```

```

i=0
j=0
n=1
C[0][0] += A[0][1]*B[1][0]
    
```



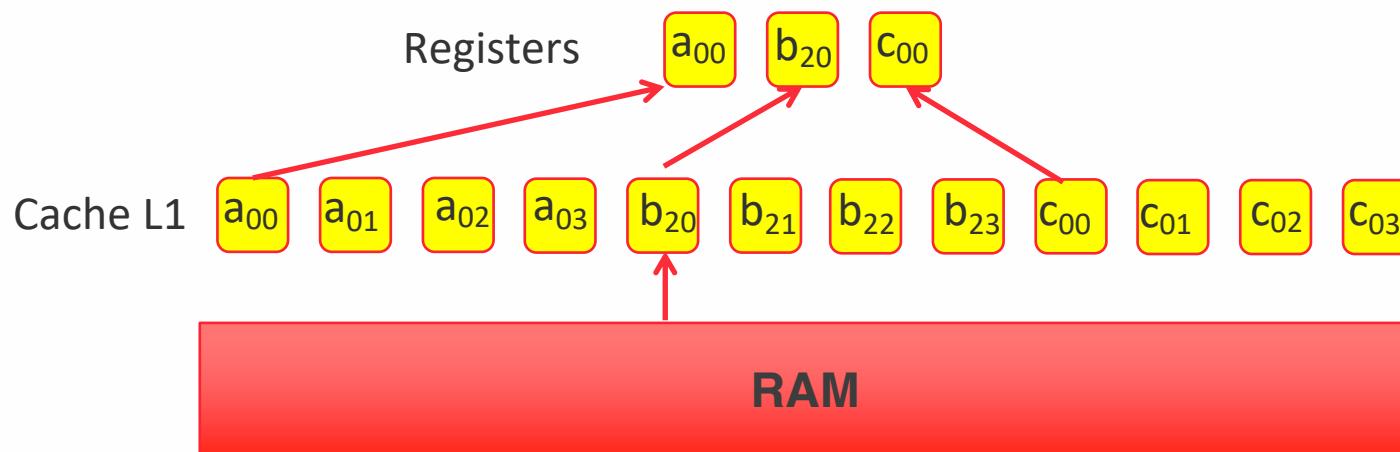


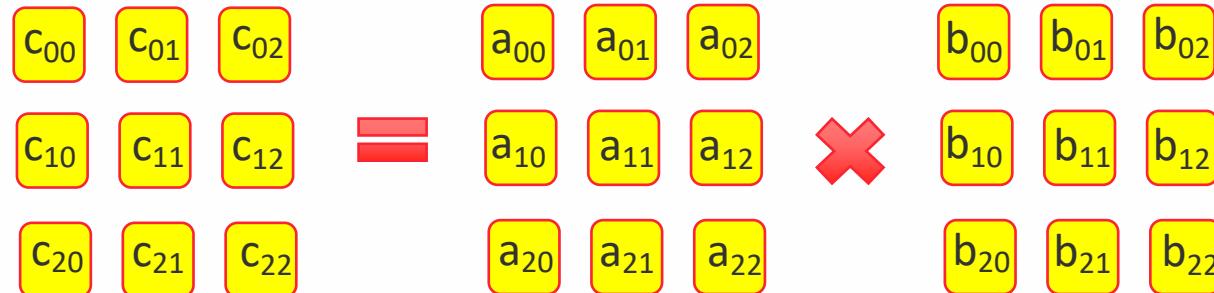
```

for i from 0 to N-1:
    for j from 0 to N-1:
        for n from 0 to N-1:
            C[i][j] += A[i][n]*B[n][j]
    
```

```

i=0
j=0
n=2
C[0][0] += A[0][2]*B[2][0]
    
```





```

for n from 0 to N-1:
    for i from 0 to N-1:
        for j from 0 to N-1:

            C[i][j] += A[i][n]*B[n][j]

```

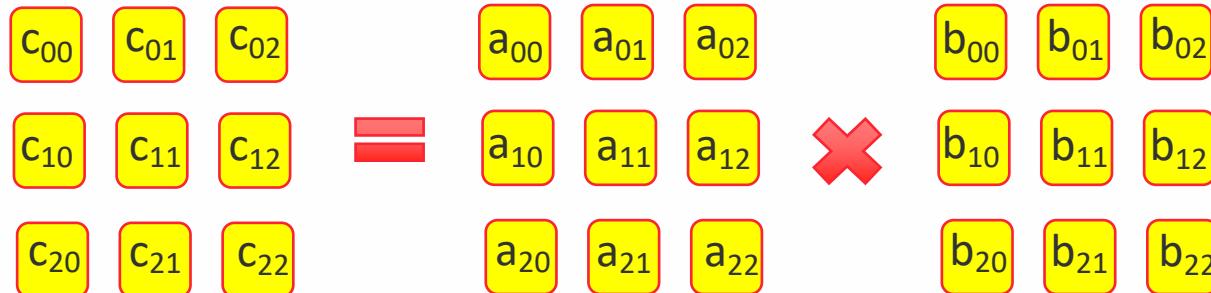
Registers



Cache L1



RAM

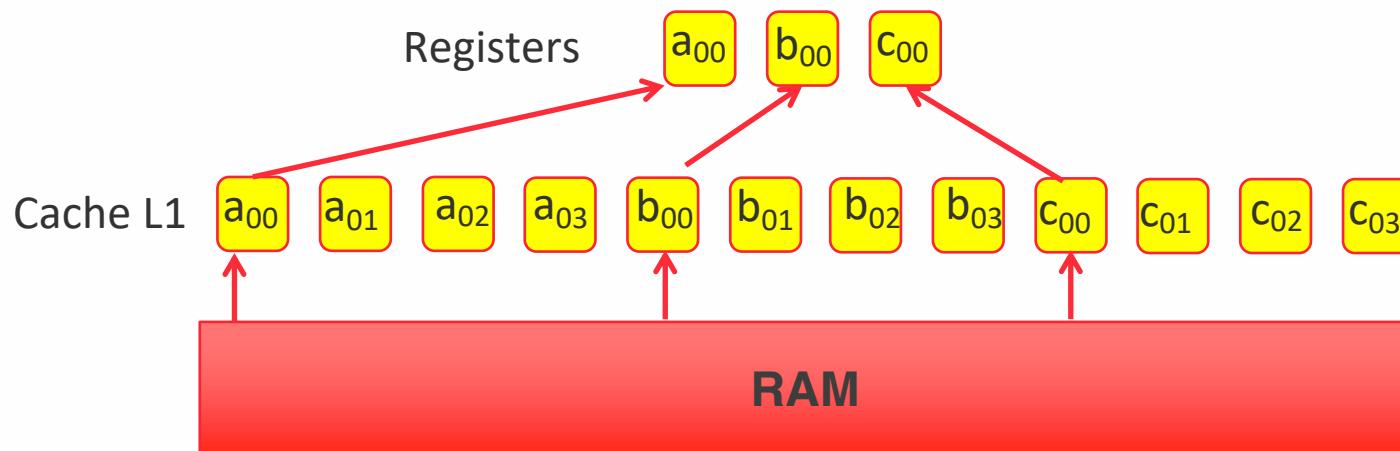


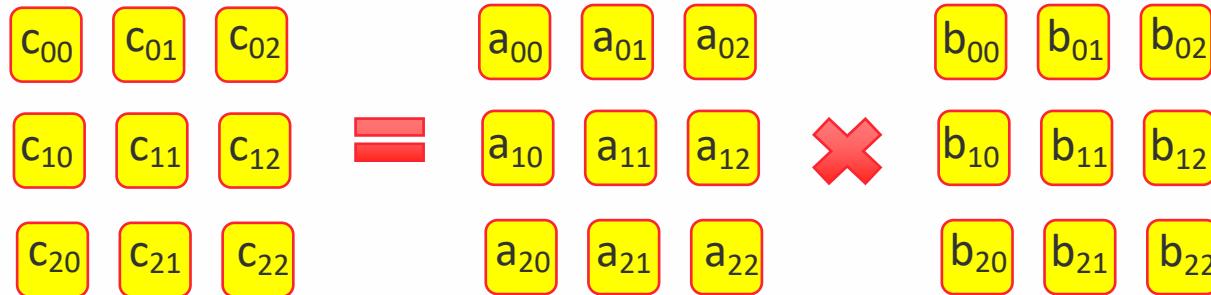
```

for n from 0 to N-1:
    for i from 0 to N-1:
        for j from 0 to N-1:
            C[i][j] += A[i][n]*B[n][j]
    
```

```

n=0
i=0
j=0
C[0][0] += A[0][0]*B[0][0]
    
```



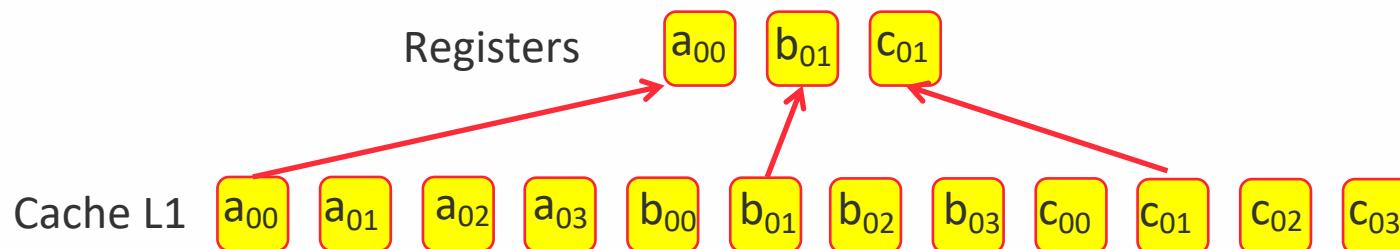


```

for n from 0 to N-1:
    for i from 0 to N-1:
        for j from 0 to N-1:
            C[i][j] += A[i][n]*B[n][j]
    
```

```

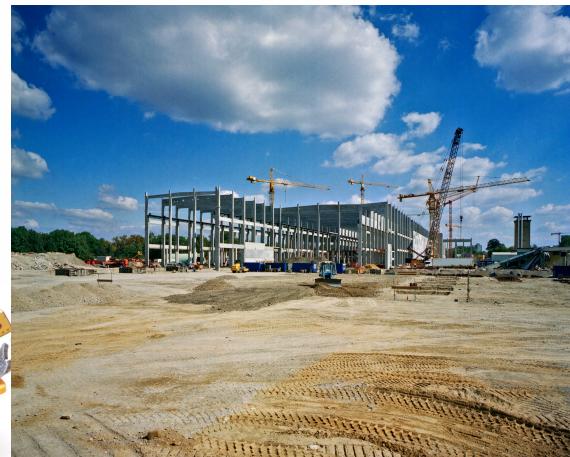
n=0
i=0
j=1
C[0][1] += A[0][0]*B[0][1]
    
```



RAM

## 2. Процессы, потоки, pthreads, конкуренция

# Processes and threads



# Processes and threads



Thread



Thread

Address space



pthreads

## POSIX

The Portable Operating System Interface is a family of standards specified by the IEEE Computer Society for maintaining compatibility between operating systems.

First published in 1988.

Among other things the committee introduced **C POSIX library** – a bunch of header files that need to work on an OS for **POSIX-compliance**.



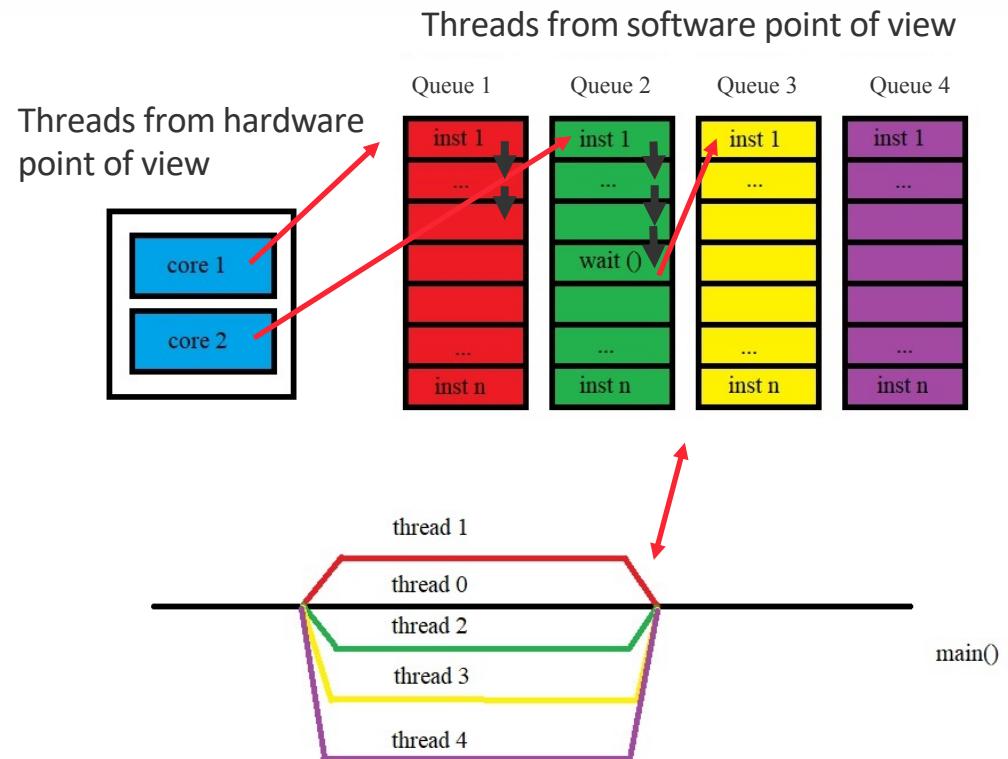
# C POSIX library

<code>&lt;aio.h&gt;</code>	Asynchronous input and output	<code>&lt;signal.h&gt;</code>	Signals, see C signal handling
<code>&lt;arpa/inet.h&gt;</code>	Functions for manipulating numeric IP addresses (part of Berkeley sockets)	<code>&lt;spawn.h&gt;</code>	Process spawning
<code>&lt;assert.h&gt;</code>	Verify assumptions	<code>&lt;stdarg.h&gt;</code>	Handle Variable Argument List
<code>&lt;complex.h&gt;</code>	Complex Arithmetic, see C mathematical functions	<code>&lt;stdbool.h&gt;</code>	Boolean type and values, see C data types
<code>&lt;cpio.h&gt;</code>	Magic numbers for the cpio archive format	<code>&lt;stddef.h&gt;</code>	Standard type definitions, see C data types
<code>&lt;ctype.h&gt;</code>	Character types	<code>&lt;stdint.h&gt;</code>	Integer types, see C data types
<code>&lt;dirent.h&gt;</code>	Allows the opening and listing of directories	<code>&lt;stdio.h&gt;</code>	Standard buffered input/output, see C file input/output
<code>&lt;dlfcn.h&gt;</code>	Dynamic linking	<code>&lt;stdlib.h&gt;</code>	Standard library definitions, see C standard library
<code>&lt;errno.h&gt;</code>	Retrieving Error Number	<code>&lt;string.h&gt;</code>	Several String Operations, see C string handling
<code>&lt;fcntl.h&gt;</code>	File opening, locking and other operations	<code>&lt;strings.h&gt;</code>	Case-insensitive string comparisons
<code>&lt;fenv.h&gt;</code>	Floating-Point Environment (FPE), see C mathematical functions	<code>&lt;stropts.h&gt;</code>	Stream manipulation, including ioctl
<code>&lt;float.h&gt;</code>	Floating-point types, see C data types	<code>&lt;sys/ipc.h&gt;</code>	Inter-process communication (IPC)
<code>&lt;fmtmsg.h&gt;</code>	Message display structures	<code>&lt;sys/mman.h&gt;</code>	Memory management, including POSIX shared memory and memory mapped files
<code>&lt;fnmatch.h&gt;</code>	Filename matching	<code>&lt;sys/msg.h&gt;</code>	POSIX message queues
<code>&lt;ftw.h&gt;</code>	File tree traversal	<code>&lt;sys/resource.h&gt;</code>	Resource usage, priorities, and limiting
<code>&lt;glob.h&gt;</code>	Pathname "globbing" (pattern-matching)	<code>&lt;sys/select.h&gt;</code>	Synchronous I/O multiplexing
<code>&lt;grp.h&gt;</code>	User group information and control	<code>&lt;sys/sem.h&gt;</code>	XSI (SysV style) semaphores
<code>&lt;iconv.h&gt;</code>	Codeset conversion facility	<code>&lt;sys/shm.h&gt;</code>	XSI (SysV style) shared memory
<code>&lt;inttypes.h&gt;</code>	Fixed sized integer types, see C data types	<code>&lt;sys/socket.h&gt;</code>	Main Berkeley sockets header
<code>&lt;iso646.h&gt;</code>	Alternative spellings, see C alternative tokens	<code>&lt;sys/stat.h&gt;</code>	File information ( <code>stat</code> et al.)
<code>&lt;langinfo.h&gt;</code>	Language information constants – builds on C localization functions	<code>&lt;sys/statvfs.h&gt;</code>	File System information
<code>&lt;libgen.h&gt;</code>	Pathname manipulation	<code>&lt;sys/time.h&gt;</code>	Time and date functions and structures
<code>&lt;limits.h&gt;</code>	Implementation-defined constants, see C data types	<code>&lt;sys/types.h&gt;</code>	Various data types used elsewhere
<code>&lt;locale.h&gt;</code>	Category macros, see C localization functions	<code>&lt;sys/uio.h&gt;</code>	Vectored I/O operations
<code>&lt;math.h&gt;</code>	Mathematical declarations, see C mathematical functions	<code>&lt;sys/un.h&gt;</code>	Unix domain sockets
<code>&lt;monetary.h&gt;</code>	String formatting of monetary units	<code>&lt;sys/utsname.h&gt;</code>	Operating system information, including <code>uname</code>
<code>&lt;nqueue.h&gt;</code>	Message queue	<code>&lt;sys/wait.h&gt;</code>	Status of terminated child processes (see <code>wait</code> )
<code>&lt;ndbm.h&gt;</code>	NDBM database operations	<code>&lt;syslog.h&gt;</code>	System error logging
<code>&lt;net/if.h&gt;</code>	Listing of local network interfaces	<code>&lt;tar.h&gt;</code>	Magic numbers for the tar archive format
<code>&lt;netdb.h&gt;</code>	Translating protocol and host names into numeric addresses (part of Berkeley sockets)	<code>&lt;termios.h&gt;</code>	Allows terminal I/O interfaces
<code>&lt;netinet/in.h&gt;</code>	Defines Internet protocol and address family (part of Berkeley sockets)	<code>&lt;tgmath.h&gt;</code>	Type-Generic Macros, see C mathematical functions
<code>&lt;netinet/tcp.h&gt;</code>	Additional TCP control options (part of Berkeley sockets)	<code>&lt;time.h&gt;</code>	Type-Generic Macros, see C date and time functions
<code>&lt;nl_types.h&gt;</code>	Localization message catalog functions	<code>&lt;trace.h&gt;</code>	Tracing of runtime behavior (DEPRECATED)
<code>&lt;poll.h&gt;</code>	Asynchronous file descriptor multiplexing	<code>&lt;ulimit.h&gt;</code>	Resource limiting (DEPRECATED in favor of <code>&lt;sys/resource.h&gt;</code> )
<code>&lt;pthread.h&gt;</code>	Defines an API for creating and manipulating POSIX threads	<code>&lt;unistd.h&gt;</code>	Various essential POSIX functions and constants
<code>&lt;pwd.h&gt;</code>	passwd (user information) access and control	<code>&lt;utime.h&gt;</code>	inode access and modification times
<code>&lt;regex.h&gt;</code>	Regular expression matching	<code>&lt;utmpx.h&gt;</code>	User accounting database functions
<code>&lt;sched.h&gt;</code>	Execution scheduling	<code>&lt;wchar.h&gt;</code>	Wide-Character Handling, see C string handling
<code>&lt;search.h&gt;</code>	Search tables	<code>&lt;wctype.h&gt;</code>	Wide-Character Classification and Mapping Utilities, see C character classification
<code>&lt;semaphore.h&gt;</code>	POSIX semaphores	<code>&lt;wordexp.h&gt;</code>	Word-expansion like the <code>shell</code> would perform
<code>&lt;setjmp.h&gt;</code>	Stack environment declarations		

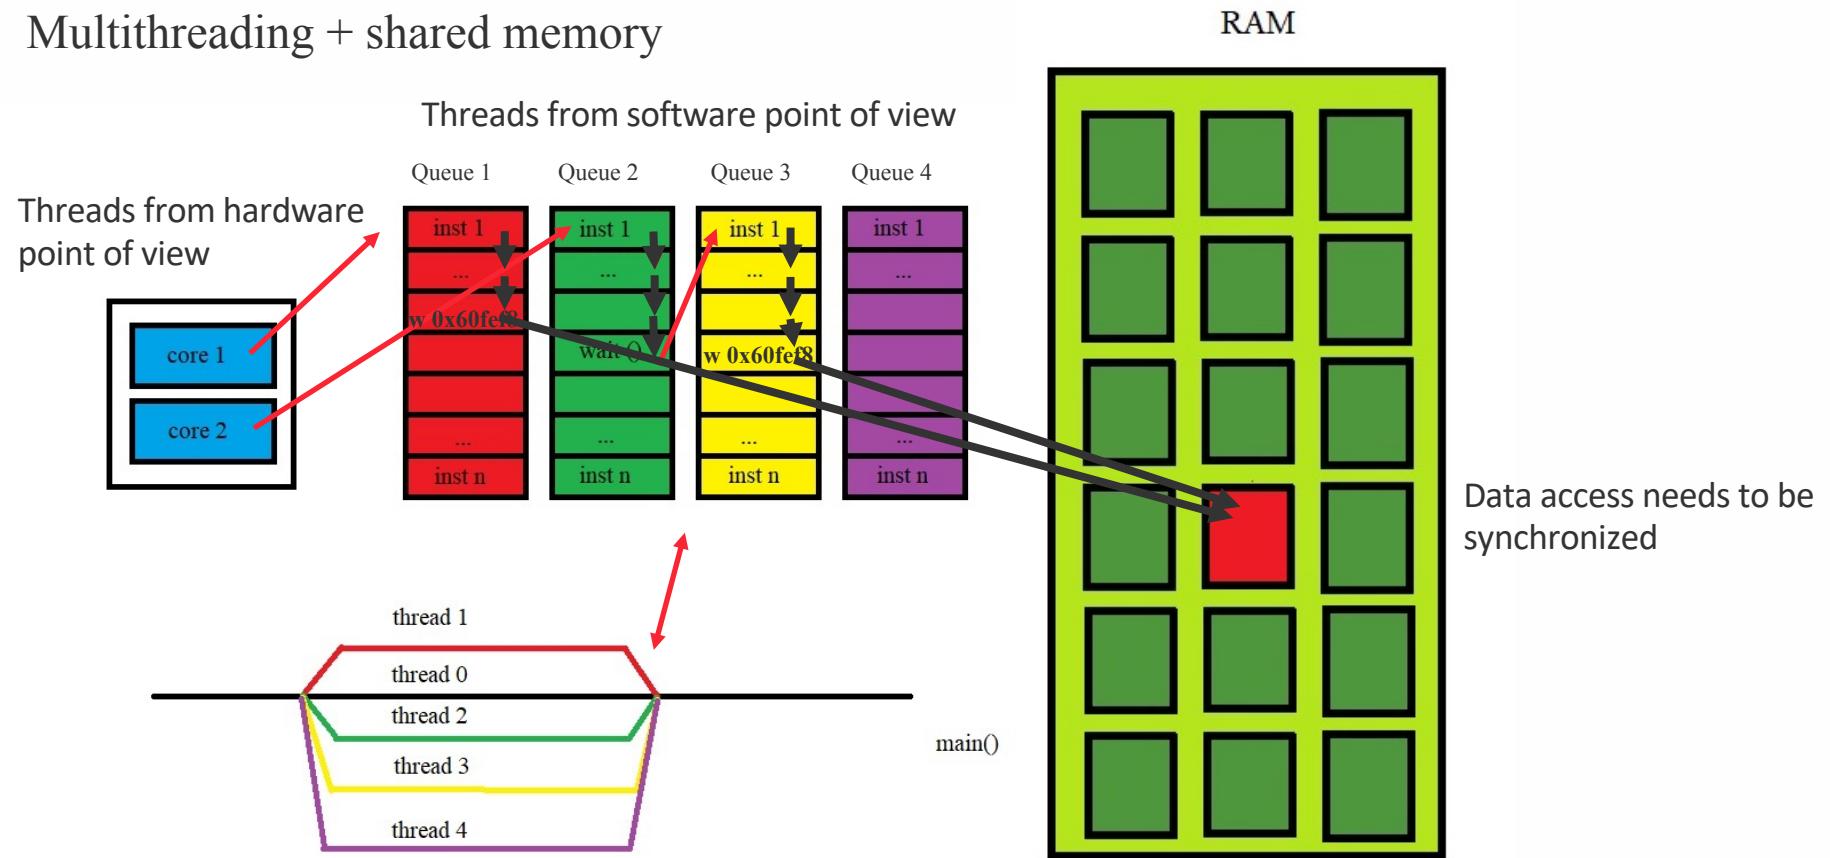
# C POSIX library (Multithreading header files)

<code>&lt;aio.h&gt;</code>	Asynchronous input and output	<code>&lt;signal.h&gt;</code>	Signals, see C signal handling
<code>&lt;arpa/inet.h&gt;</code>	Functions for manipulating numeric IP addresses (part of Berkeley sockets)	<code>&lt;spawn.h&gt;</code>	Process spawning
<code>&lt;assert.h&gt;</code>	Verify assumptions	<code>&lt;stdarg.h&gt;</code>	Handle Variable Argument List
<code>&lt;complex.h&gt;</code>	Complex Arithmetic, see C mathematical functions	<code>&lt;stdbool.h&gt;</code>	Boolean type and values, see C data types
<code>&lt;cpio.h&gt;</code>	Magic numbers for the cpio archive format	<code>&lt;stddef.h&gt;</code>	Standard type definitions, see C data types
<code>&lt;ctype.h&gt;</code>	Character types	<code>&lt;stdint.h&gt;</code>	Integer types, see C data types
<code>&lt;dirent.h&gt;</code>	Allows the opening and listing of directories	<code>&lt;stdio.h&gt;</code>	Standard buffered input/output, see C file input/output
<code>&lt;dlfcn.h&gt;</code>	Dynamic linking	<code>&lt;stdlib.h&gt;</code>	Standard library definitions, see C standard library
<code>&lt;errno.h&gt;</code>	Retrieving Error Number	<code>&lt;string.h&gt;</code>	Several String Operations, see C string handling
<code>&lt;fcntl.h&gt;</code>	File opening, locking and other operations	<code>&lt;strings.h&gt;</code>	Case-insensitive string comparisons
<code>&lt;fenv.h&gt;</code>	Floating-Point Environment (FPE), see C mathematical functions	<code>&lt;stropts.h&gt;</code>	Stream manipulation, including ioctl
<code>&lt;float.h&gt;</code>	Floating-point types, see C data types	<code>&lt;sys/ipc.h&gt;</code>	Inter-process communication (IPC)
<code>&lt;fmtmsg.h&gt;</code>	Message display structures	<code>&lt;sys/mman.h&gt;</code>	Memory management, including POSIX shared memory and memory mapped files
<code>&lt;fnmatch.h&gt;</code>	Filename matching	<code>&lt;sys/msg.h&gt;</code>	POSIX message queues
<code>&lt;ftw.h&gt;</code>	File tree traversal	<code>&lt;sys/resource.h&gt;</code>	Resource usage, priorities, and limiting
<code>&lt;glob.h&gt;</code>	Pathname "globbing" (pattern-matching)	<code>&lt;sys/select.h&gt;</code>	Synchronous I/O multiplexing
<code>&lt;grp.h&gt;</code>	User group information and control	<code>&lt;sys/sem.h&gt;</code>	XSI (SysV style) semaphores
<code>&lt;iconv.h&gt;</code>	Codeset conversion facility	<code>&lt;sys/shm.h&gt;</code>	XSI (SysV style) shared memory
<code>&lt;inttypes.h&gt;</code>	Fixed sized integer types, see C data types	<code>&lt;sys/socket.h&gt;</code>	Main Berkeley sockets header
<code>&lt;iso646.h&gt;</code>	Alternative spellings, see C alternative tokens	<code>&lt;sys/stat.h&gt;</code>	File information ( <code>stat</code> et al.)
<code>&lt;langinfo.h&gt;</code>	Language information constants – builds on C localization functions	<code>&lt;sys/statvfs.h&gt;</code>	File System information
<code>&lt;libgen.h&gt;</code>	Pathname manipulation	<code>&lt;sys/time.h&gt;</code>	Time and date functions and structures
<code>&lt;limits.h&gt;</code>	Implementation-defined constants, see C data types	<code>&lt;sys/types.h&gt;</code>	File access and modification times
<code>&lt;locale.h&gt;</code>	Category macros, see C localization functions	<code>&lt;sys/uio.h&gt;</code>	Various data types used elsewhere
<code>&lt;math.h&gt;</code>	Mathematical declarations, see C mathematical functions	<code>&lt;sys/un.h&gt;</code>	Vectored I/O operations
<code>&lt;monetary.h&gt;</code>	String formatting of monetary units	<code>&lt;sys/utsname.h&gt;</code>	Unix domain sockets
<code>&lt;mqqueue.h&gt;</code>	Message queue	<code>&lt;sys/wait.h&gt;</code>	Operating system information, including <code>uname</code>
<code>&lt;ndbm.h&gt;</code>	NDBM database operations	<code>&lt;syslog.h&gt;</code>	Status of terminated child processes (see <code>wait</code> )
<code>&lt;net/if.h&gt;</code>	Listing of local network interfaces	<code>&lt;tar.h&gt;</code>	System error logging
<code>&lt;netdb.h&gt;</code>	Translating protocol and host names into numeric addresses (part of Berkeley sockets)	<code>&lt;termios.h&gt;</code>	Magic numbers for the tar archive format
<code>&lt;netinet/in.h&gt;</code>	Defines Internet protocol and address family (part of Berkeley sockets)	<code>&lt;tgmath.h&gt;</code>	Allows terminal I/O interfaces
<code>&lt;netinet/tcp.h&gt;</code>	Additional TCP control options (part of Berkeley sockets)	<code>&lt;time.h&gt;</code>	Type-Generic Macros, see C mathematical functions
<code>&lt;nl_types.h&gt;</code>	Localization message catalog functions	<code>&lt;trace.h&gt;</code>	Tracing of runtime behavior (DEPRECATED)
<code>&lt;poll.h&gt;</code>	Asynchronous file descriptor multiplexing	<code>&lt;ulimit.h&gt;</code>	Resource limiting (DEPRECATED in favor of <code>&lt;sys/resource.h&gt;</code> )
<code>&lt;pthread.h&gt;</code>	Defines an API for creating and manipulating POSIX threads	<code>&lt;unistd.h&gt;</code>	Various essential POSIX functions and constants
<code>&lt;pwd.h&gt;</code>	passwd (user information) access and control	<code>&lt;utime.h&gt;</code>	inode access and modification times
<code>&lt;regex.h&gt;</code>	Regular expression matching	<code>&lt;utmpx.h&gt;</code>	User accounting database functions
<code>&lt;sched.h&gt;</code>	Execution scheduling	<code>&lt;wchar.h&gt;</code>	Wide-Character Handling, see C string handling
<code>&lt;search.h&gt;</code>	Search tables	<code>&lt;wctype.h&gt;</code>	Wide-Character Classification and Mapping Utilities, see C character classification
<code>&lt;semaphore.h&gt;</code>	POSIX semaphores	<code>&lt;wordexp.h&gt;</code>	Word-expansion like the <code>shell</code> would perform
<code>&lt;setjmp.h&gt;</code>	Stack environment declarations		

# Multithreading



## Multithreading + shared memory



## A short example

```
int counter;
pthread_mutex_t lock;

void* doSomeThing(void *arg)
{
    pthread_mutex_lock(&lock);
    counter += 1;
    pthread_mutex_unlock(&lock);
    return NULL;
}

int main(void)
{
    int N = 500;
    pthread_t tid[N];
    int i = 0;
    int err;
    pthread_mutex_init(&lock, NULL);
    for(i=0; i<N; i++) pthread_create(&(tid[i]), NULL, &doSomeThing, NULL);
    for(i=0; i<N; i++) pthread_join(tid[i], NULL);
    pthread_mutex_destroy(&lock);
    return 0;
}
```

Global variables. They become automatically shared between threads.

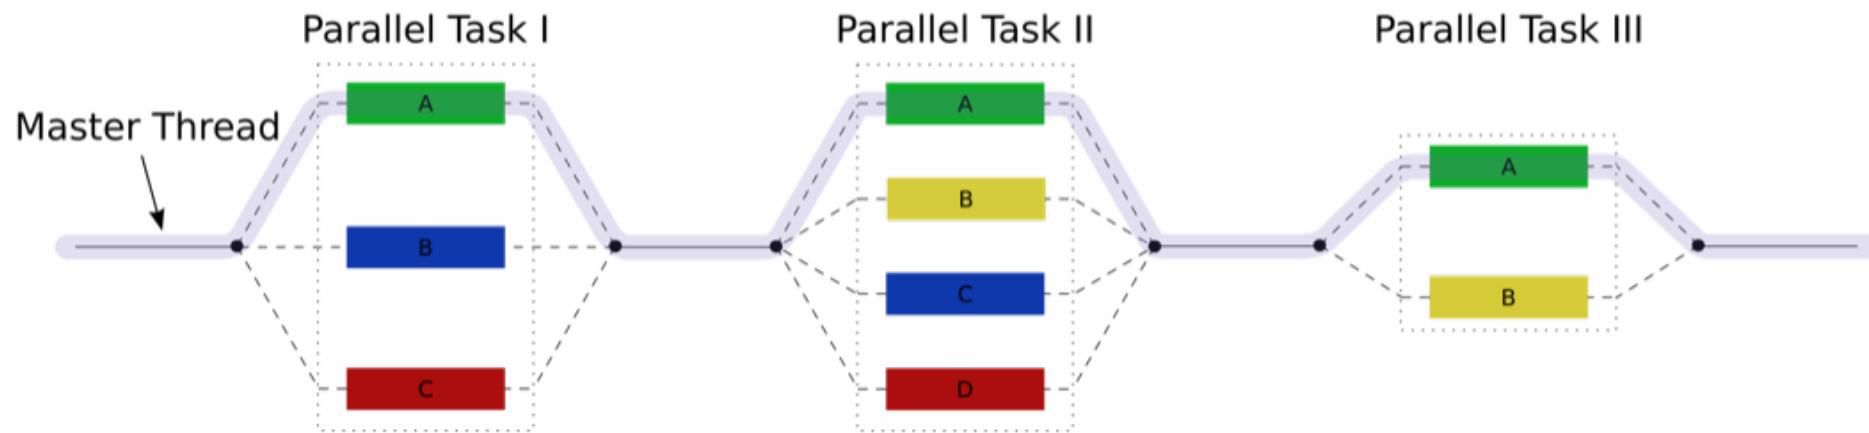
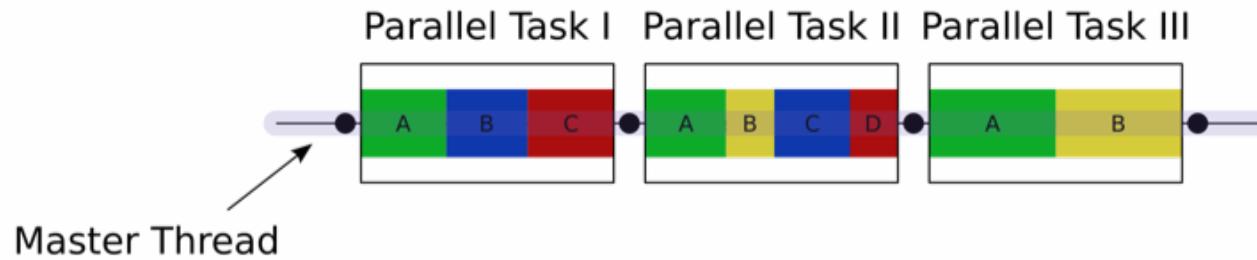
Notice the definition. The function has to have 1 argument, which is pointer to void and return pointer to void. Before using *arg* cast it to the needed type. For ex: *(int) arg*.

You can also write *pthread\_exit(0)*; instead of *return NULL;*

4th argument is arguments that are passed to *doSomeThing*. We can pass there whatever we want but we need to cast it as poniter to void. For ex: *(void\*) i*.

Compile command looks something like this: **gcc prog.c -o prog.exe -pthread -lm**

# Fork-join concept



# Race condition (состояние гонки)

Example: calculate the sum of array elements  
“REDUCTION(SUM)”



$s = \text{sum}(a[i])$

$s=s+a[i]$



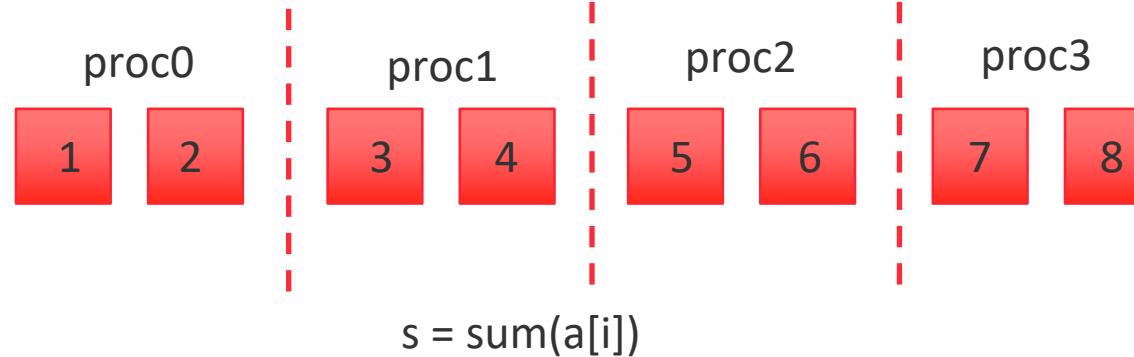
1. Read  $s$  from memory
2. Add a value
3. Write the updated value back



memory element that holds the variable 's' which is the sum of the elements

# Race condition (состояние гонки)

Example: calculate the sum of array elements  
“REDUCTION(SUM)”



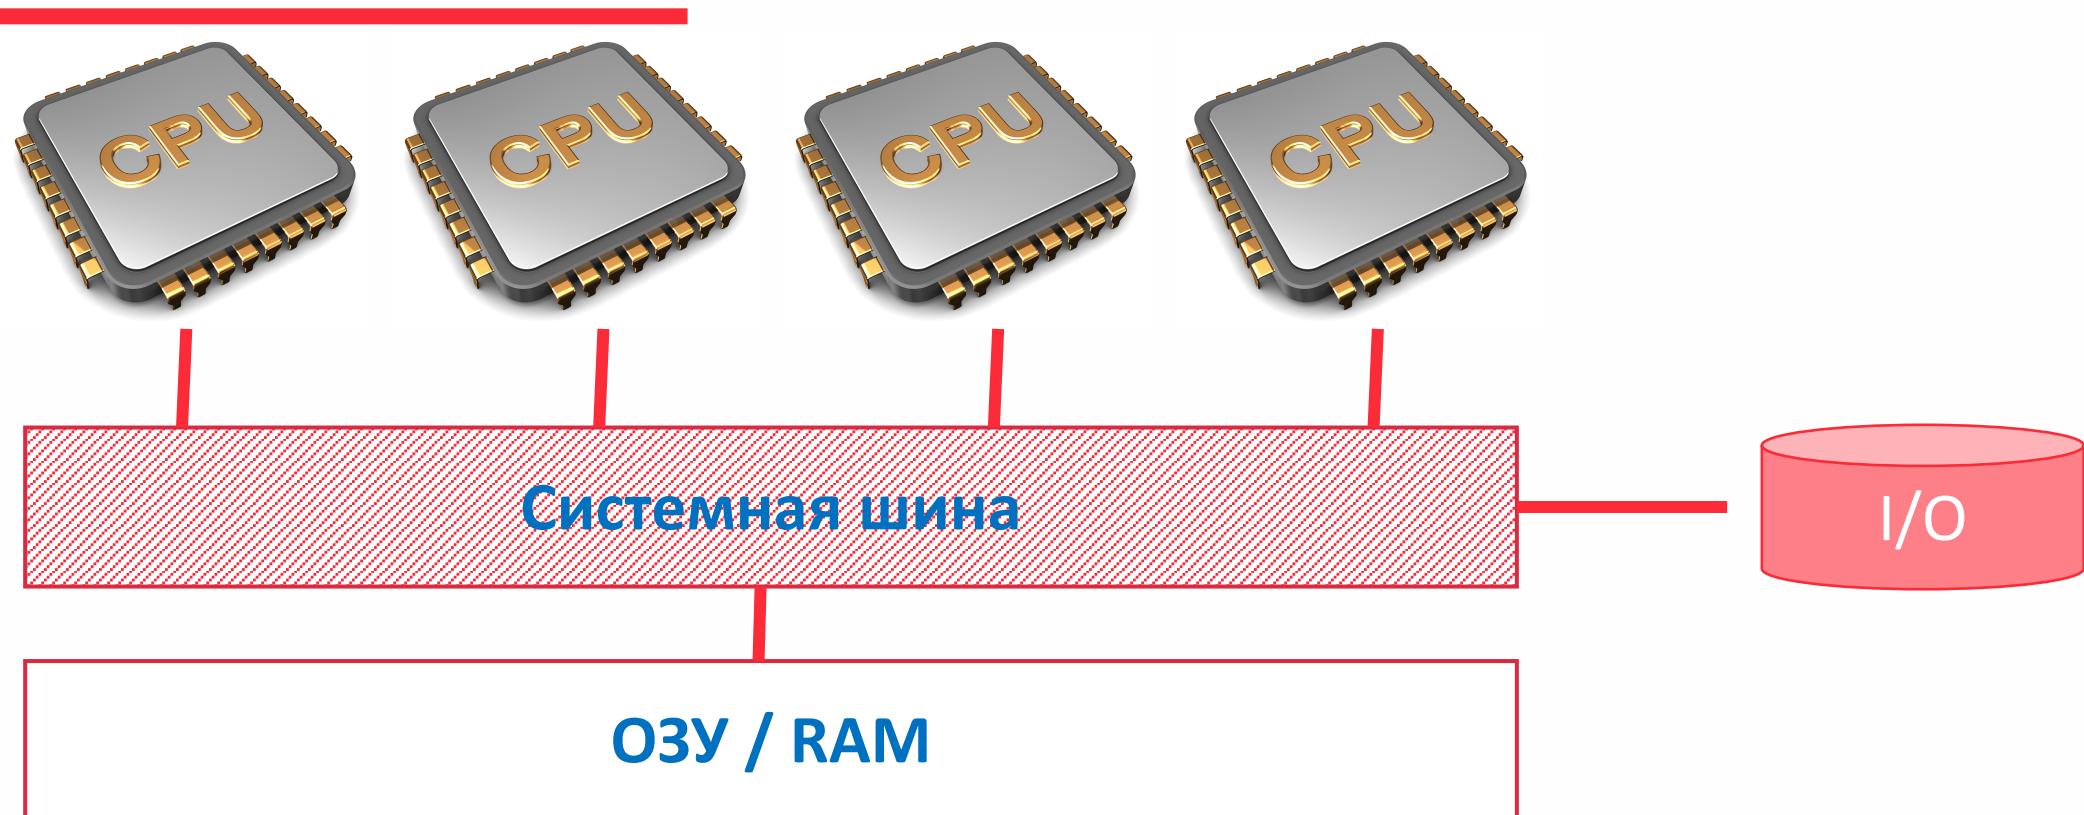
- $s=s+a[i]$  →
1. Read  $s$  from memory
  2. Add a value
  3. Write the updated value back



memory element that holds the variable 's' which is the sum of the elements

### 3. OpenMP framework

# OpenMP. Системы с общей памятью



OpenMP:

- API для написания многопоточных приложений
- Набор директив компилятора, функций и переменных окружения
- Значительно упрощает разработку многопоточных приложений (особенно в HPC)



# OpenMP. Директивы компилятора

---

В большей степени OpenMP – это директивы компилятора. В общем виде:

```
#pragma omp construct [clause [clause]...]
#pragma omp parallel num_threads(4)
```

Прототипы функций и типы данных находятся в header-файле:

```
#include <omp.h>
```

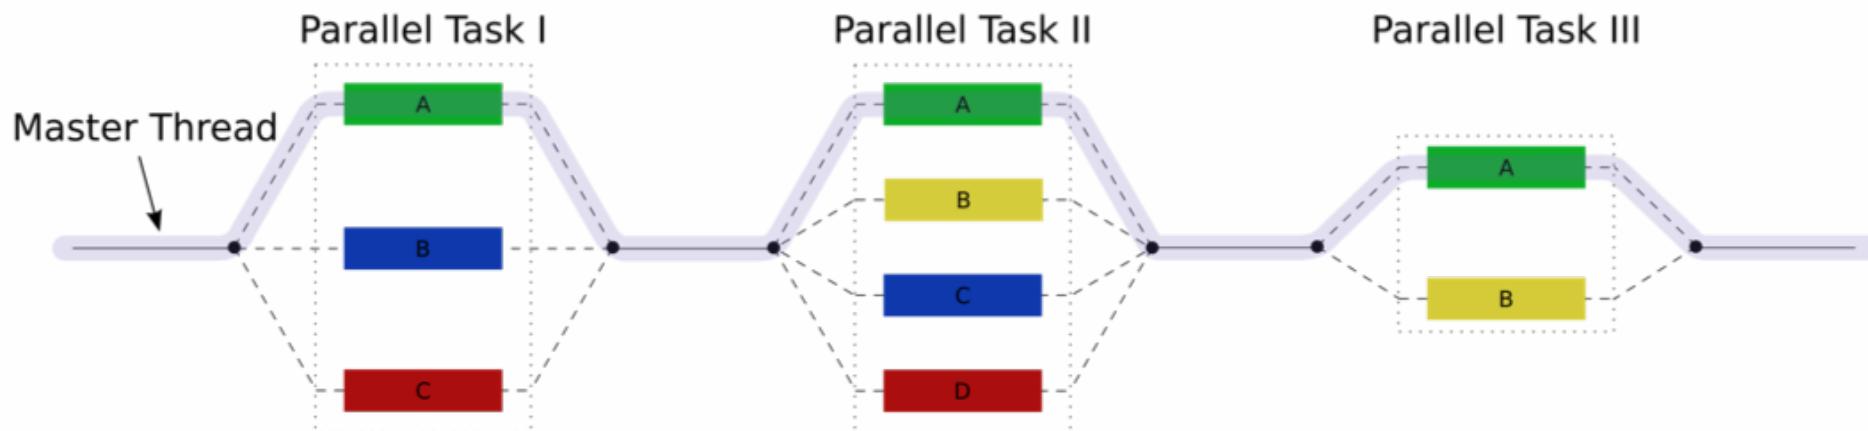
Переменные окружения можно задать на системном уровне, например:

*OMP\_NUM\_THREADS*

# OpenMP. Концепция Fork-Join

У разработчика есть явный контроль параллелизации за счет модели параллелизма типа fork-join:

- все OpenMP программы начинаются как один процесс (мастер-поток, master thread), который выполняется до директивы компилятора (parallel region construct)
- **FORK**: мастер-поток создает команду параллельных потоков
- **JOIN**: когда потоки заканчивают выполнение параллельных инструкций, они синхронизуются и завершают свою жизнь, остается только мастер-поток.





# OpenMP. Концепция Fork-Join

---

**Разработчик** вставляет в код директивы, которые поясняют компилятору:

- какие части программы должны исполняться параллельно
- какие задания будут выполнять различные потоки
- какие данные общие (shared), какие данные приватные (private)
- `#pragma omp`

**Компилятор** генерирует многопоточный код

**Негласное правило:** один поток на ядро (2 или 4 если используется НТ)

Работа с данными в параллельных частях программы приводит к необходимости **синхронизации потоков**.

# OpenMP. Hello, world!

---

```
#include <omp.h>
#include <stdio.h>

int main() {

#pragma omp parallel
printf("Hello world from thread %d\n",omp_get_thread_num());

return 0;
}
```

**Compile line:**

```
gcc -fopenmp helloWorld.c
```



# OpenMP. Как задать кол-во потоков

---

**Через переменную окружения:**

```
setenv OMP_NUM_THREADS 2
```

(cshel1)

```
export OMP_NUM_THREADS=2
```

(bash)

**Функция библиотеки OpenMP:**

```
omp_set_num_threads(2);
```

**В директиве компилятору:**

```
#pragma omp parallel num_threads(2)
```



# OpenMP. Конструкция parallel

---

```
#include <omp.h>

int main() {
    int var1, var2, var3;

    ...serial Code

    #pragma omp parallel private(var1, var2) shared (var3)
    {
        ...parallel section
    }

    ...resume serial code

}
```



# OpenMP. Конструкция parallel

---

- Когда поток достигает директивы `parallel`, он становится мастер-потоком и его номер становится равен 0 (остальные потоки имеют уникальные номера)
- Оптимально если все потоки исполняют один и тот же код, но можно делать параллелизм по задачам (`task parallelism, work sharing`)
- В конце конструкции `parallel` всегда **барьер!** Только мастер-поток продолжает исполнение программы
- Если какой-то поток аварийно останавливается (`terminate`) внутри параллельной части, то все потоки останавливаются и результат не определен.
- Если программа скомпилирована без использования OpenMP – она компилируется (возможно, за исключением части функций) и работает последовательно.
- Простота параллелизации!
- Простота поддержки кода (`code maintenance`) – одна версия для параллельной и последовательной программы.

```
static void critical_example(int v) {

    // Shared variables
    int a = 0;
    int b = v;

    #pragma omp parallel
    {
        // Private variable - one for each thread
        int c;

        // Reading from "b" is safe: it is read-only
        // Writing to "c" is safe: it is private
        c = b;

        // Reading from "c" is safe: it is private
        // Writing to "c" is safe: it is private
        c = c * 10;

        #pragma omp critical
        {
            // Any modifications of "a" are safe:
            // we are inside a critical section
            a = a + c;
        }
    }

    // Reading from "a" is safe:
    // we are outside parallel region
    std::cout << a << std::endl;
}
```

shared  
private  
critical

# OpenMP: Небольшое сравнение с pthreads

```
void thunk ()  
{  
    foobar ();  
}  
  
pthread_t tid[4];  
for (int i = 1; i < 4; ++i)  
    pthread_create (  
        &tid[i], 0, thunk, 0);  
think();  
  
for (int i = 1; i < 4; ++i)  
    pthread_join (tid[i]);
```



```
#pragma omp parallel num_threads(4)  
{  
    foobar ();  
}
```



# OpenMP. Work-sharing constructs

---

- Не создает новых потоков, только разделяет работу между всеми потоками в команде
- **На входе нет барьера, на выходе есть!** (но можно отменить барьер на выходе)

3 типа work-sharing constructs:

- **for loop:** параллелизация циклов (data parallelism)
- **sections:** разделение на параллельные участки, в каждом участке свои инструкции
- **single:** часть кода, которая исполняется строго одним потоком
- **master:** то же, что и single, но выполняет только поток с номером 0 (мастер)



# OpenMP. Work-sharing constructs

---

- Не создает новых потоков, только разделяет работу между всеми потоками в команде
- **На входе нет барьера, на выходе есть!** (но можно отменить барьер на выходе)

3 типа work-sharing constructs:

- **for loop**: параллелизация циклов (data parallelism)
- **sections**: разделение на параллельные участки, в каждом участке свои инструкции
- **single**: часть кода, которая исполняется строго одним потоком
- **master**: то же, что и single, но выполняет только поток с номером 0 (мастер)

# OpenMP. Пример parallel for

---

```
int main()
{
    int sum=0;
    int a[N], i;

#pragma omp parallel for
for(i=0;i<N;i++)
{
    a[i]=i;
    printf("iterate i=%3d by thread %3d\n",i, omp_get_thread_num());
}
return 0;
}
```

# OpenMP. Пример parallel for

```
int main()
{
    int sum=0;
    int a[N], i;
    #pragma omp parallel for
    for(i=0;i<N;i++)
    {
        a[i]=i;
        printf("iterate i=%3d by thread %3d\n",i, omp_get_thread_num());
    }
    return 0;
}
```

#pragma omp parallel

{

#pragma omp for



# OpenMP. for или не for?

---

```
#pragma omp parallel for
for(i=0;i<N;i++)
{
    a[i]=i;
    printf("iterate i=%3d by thread %3d\n", i, omp_get_thread_num());
}
```

---

```
#pragma omp parallel
for(i=0;i<N;i++)
{
    a[i]=i;
    printf("iterate i=%3d by thread %3d\n", i, omp_get_thread_num());
}
```

# OpenMP. Общий вид директив

```
#pragma omp directive [clause ...]
                      if (scalar_expression)
                      private (list)
                      shared (list)
                      default (shared | none)
                      firstprivate (list)
                      reduction (operator: list)
                      copyin (list)
                      num_threads (integer-expression)
```

- case sensitive
- одна директива в конструкции
- если следуют больше одной инструкции, надо обернуть в {}
- длинную строку можно переносить обратным слэшем (“\\”)

# OpenMP. parallel for

---

```
#pragma omp for [clause ...]
    schedule (type [,chunk])
    private (list)
    firstprivate(list)
    lastprivate(list)
    shared (list)
    reduction (operator: list)
    nowait
```

**SCHEDULE:** описывает как разделять между потоками итерации цикла:

- static = статично разделять на части размера chunk, по умолчанию chunk примерно одинаковый для всех потоков
- dynamic = разделено на части chunk, но кто первый пришел, тот и забрал очередной chunk.
- guided = размер chunk уменьшается со временем
- runtime = определяется переменной окружения OMP\_SCHEDULE

# OpenMP. Пример parallel for

```
#pragma omp parallel shared(n,a,b,x,y), private(i)
{ // start parallel region

    #pragma omp for nowait
    for (i=0;i<n;i++)
        b[i] = += a[i];

    #pragma omp for nowait
    for (i=0;i<n;i++)
        x[i] = 1./y[i];

} // end parallel region (implied barrier)
```

**fork** занимает время – не выходите без надобности из конструкции **parallel**  
**nowait** - отменяет барьер на выходе из work-sharing конструкции



# OpenMP. Пример parallel sections

---

```
#pragma omp sections [clause ...]
    private (list)
    firstprivate(list)
    lastprivate(list)
    reduction (operator: list)
    nowait
{
    #pragma omp section
    structured block
    #pragma omp section
    structured block
}
```

**nowait** - отменяет барьер на выходе из work-sharing конструкции  
возможна расходимость потоков (thread divergence)  
возможно неравномерное распределение работы



# OpenMP. Пример single

---

```
#pragma omp single [clause ...]
    private (list)
    firstprivate(list)
    nowait
structured block
```

- только один поток выполняет конструкцию `single` (кто первый пришел)
- важно если часть кода не `thread safe` (например, I/O)
- в конце барьер, но его можно отменить (`nowait`)



# OpenMP. Пример firstprivate

---

Что здесь не так?

```
#pragma omp parallel for
for (i=0;i<n;i++) {
    x = i*dx
    y(i) = exp(x) * cos(x) * sqrt(6*x+5);
}
```

# OpenMP. Пример firstprivate

---

Что здесь не так?

```
#pragma omp parallel for
for (i=0; i<n; i++) {
    x = i*dx
    y(i) = exp(x) * cos(x) * sqrt(6*x+5);
}
```

By default, `x` is shared variable (`i` is private).

Could have: Thread 0 set `x` for some `i`.

Thread 1 sets `x` for different `i`.

Thread 0 uses `x` but it is now incorrect.

# OpenMP. Пример firstprivate

---

Что здесь не так?

```
#pragma omp parallel for private(x)
for (i=0;i<n;i++) {
    x = i*dx
    y(i) = exp(x) * cos(x) * sqrt(6*x+5);
}
```

Что насчет i, dx, y?

By default dx, n, y shared.

dx, n used but not changed. y changed, but independently for each i



# OpenMP. Пример firstprivate

---

Что здесь не так?

```
dx = 1/n.;  
#pragma omp parallel for private(x,dx)  
for (i=0;i<n;i++) {  
    x = i*dx  
    y(i) = exp(x) * cos(x) * sqrt(6*x+5);  
}
```

# OpenMP. Пример firstprivate

---

Что здесь не так?

```
dx = 1/n.;  
#pragma omp parallel for private(x, dx)  
for (i=0;i<n;i++){  
    x = i*dx  
    y(i) = exp(x) * cos(x) * sqrt(6*x+5);  
}
```

private(dx) создает новую приватную переменную для каждого потока, но не гарантирует инициализацию

**firstprivate** clause creates private variables and initializes to the value from the master thread before the loop.

**lastprivate** copies last value computed by a thread (for  $i=n$ ) to the master thread copy to continue execution.

# OpenMP. clauses

These clauses not strictly necessary but may be convenient (and may have performance penalties too).

- **lastprivate** private data is undefined after parallel construct. this gives it the value of last iteration (as if sequential) or sections construct (in lexical order).
- **firstprivate** pre-initialize private vars with value of variable with same name before parallel construct.
- **default** (none | shared). In fortran can also have private. Then only need to list exceptions. (none is better habit).
- **nowait** suppress implicit barrier at end of work sharing construct. Cannot ignore at end of parallel region. (But no guarantee that if have 2 `for` loops where second depends on data from first that same threads execute same iterates)

# More clauses

- if (logical expr) true = execute parallel region with team of threads; false = run serially (loop too small, too much overhead)
- reduction for assoc. and commutative operators compiler helps out; reduction variable is shared by default (no need to specify).

```
#pragma omp parallel for default(none) \
    shared(n, a) \
    reduction(+:sum)

for (i=0;i<n;i++)
    sum += a[i]
/* end of parallel reduction */
```

Also other arithmetic and logical ops., min,max intrinsics in Fortan only.



# Потоковая безопасность (thread safety)

---

- `rand()` is not thread-safe, can lead to [Heisenbug](#) (Schrödinbug)
- `rand_r` is thread safe because the function is entirely pure.
- Thread **safe** list of functions [here](#)
- Functions **unsafe** to call the function concurrently :
  1. `Rand()`
  2. `Strtok()`
  3. `Basename()`
  4. Others [here](#)

# Генерация случайных чисел

---

```
omp_picalc.c
#pragma omp parallel ...
{
    unsigned int seed =
        123456789 *
        thread_id;
    ...
    double x =
        ((double)
        rand_r(&seed))...
}
```

## TIMING

```
> gcc omp_picalc.c -fopenmp
> time -p a.out 100000000
npoints: 100000000
hits: 78541717
pi_est: 3.141669
```

```
omp_picalc_rand_r.c:
unsigned int seed =
    123456789;
#pragma omp parallel...
{
    ...
    double x =
        ((double) rand_r(&seed))...
}
```

## TIMING

```
> gcc omp_picalc_rand_r.c -fopenmp
> time -p a.out 100000000
npoints: 100000000
hits: 77951102
pi_est: 3.118044
```

```
/*** OMP ***/
const int nThreads = 4; // number of threads to use
unsigned int seeds[nThreads];

void seedThreads() {
    int my_thread_id;
    unsigned int seed;
#pragma omp parallel private (seed, my_thread_id)
{
    my_thread_id = omp_get_thread_num();

    //create seed on thread using current time
    unsigned int seed = (unsigned) time(NULL);

    //munge the seed using our thread number so that each thread has its
    //own unique seed, therefore ensuring it will generate a different set of numbers
    seeds[my_thread_id] = (seed & 0xFFFFFFFF0) | (my_thread_id + 1);

    printf("Thread %d has seed %u\n", my_thread_id, seeds[my_thread_id]);
}

/*** OMP ***/

```

```
/*** OMP ***/
omp_set_num_threads(nThreads);
seedThreads();
/*** OMP ***/

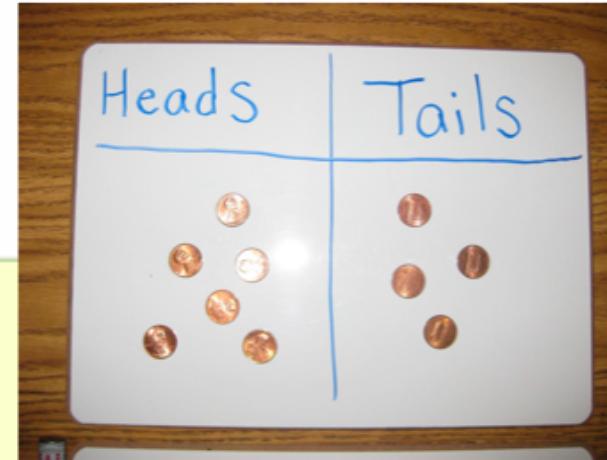
```

[Link to the file](#)

# Пример использования массива seed

```
#pragma omp parallel num_threads(nThreads) default(None) \
    private(numFlips, tid, seed) \
    shared(trialFlips, seeds) \
    reduction(+:numHeads, numTails)
{
    tid = omp_get_thread_num();      // my thread id
    seed = seeds[tid];            // it is much faster to keep a private copy of our seed
    srand(seed);                  //seed rand_r or rand

    #pragma omp for
    for (numFlips=0; numFlips<trialFlips; numFlips++) {
        // in Windows, can use rand()
        if (rand()%2 == 0) // if random number is even, call it heads
        // linux: rand_r() is thread safe, to be run on separate threads concurrently
        if (rand_r(&seed)%2 == 0) // if random number is even, call it heads
            numHeads++;
        else
            numTails++;
    }
}
```



Note: in Linux we need to use the `rand_r()` function for thread-safe generation of the numbers. However, in Windows, the `rand()` function is thread-safe.

# for loop scheduling

## Loop Scheduling - 4 Types

### Static

- ▶ So far only done static scheduling with fixed size chunks
- ▶ Threads get fixed size chunks in rotating fashion
- ▶ Great if each iteration has same work load

### Guided

- ▶ Hybrid between static/dynamic, start with each thread taking a "big" chunk
- ▶ When a thread finishes, requests a "smaller" chunk, next request is smaller

### Dynamic

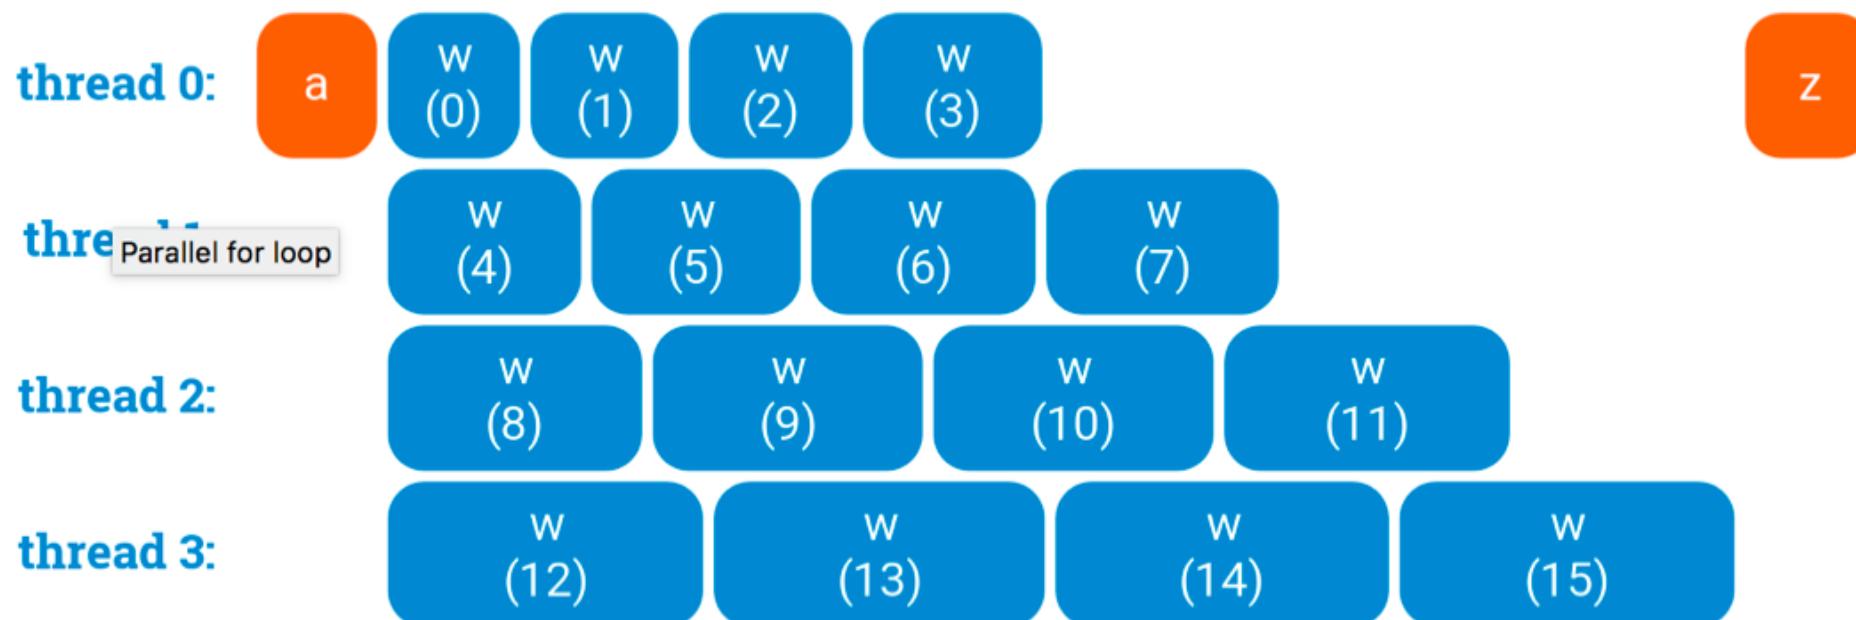
- ▶ Threads get fixed chunks but when done, request another chunk
- ▶ Incurs more overhead but balances uneven load better

### Runtime

- ▶ Environment variables used to select one of the others
- ▶ Flexible but requires user awareness: *What's an environment variable?*

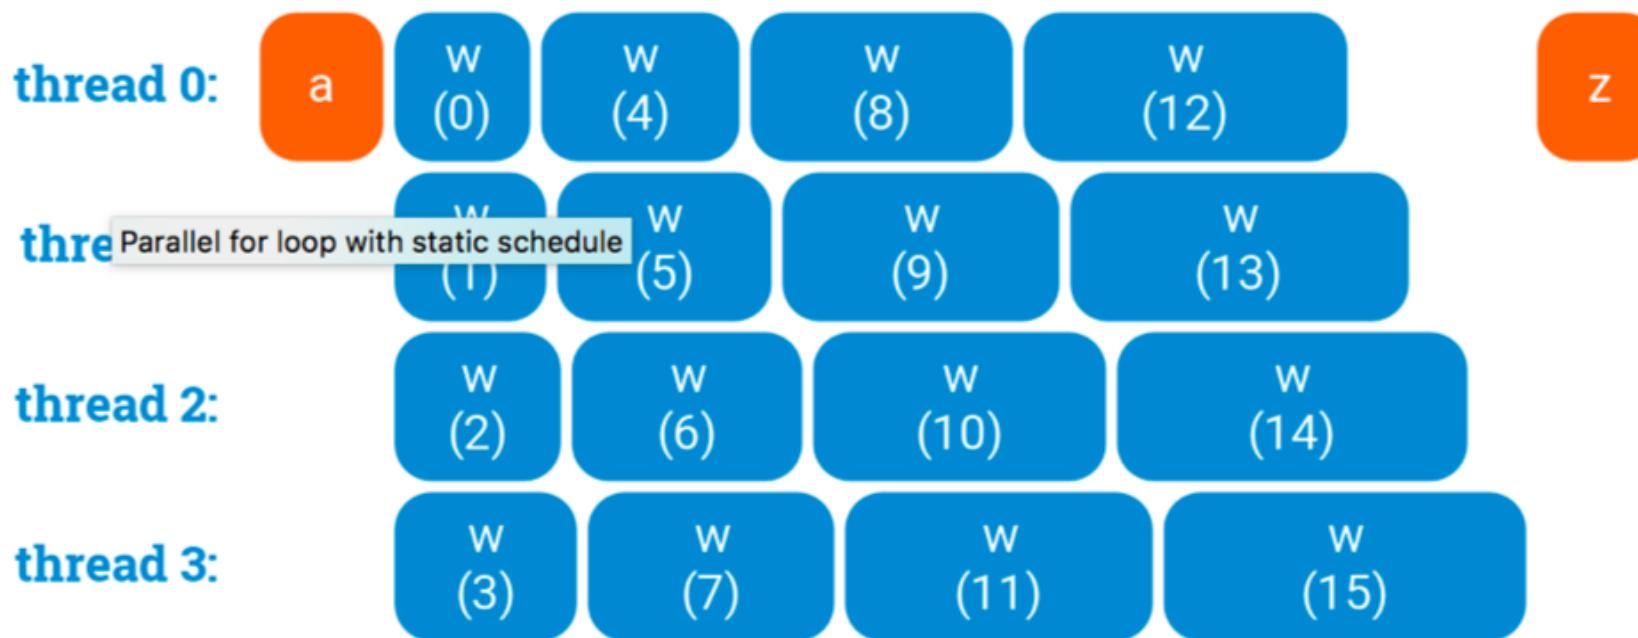
# for loop scheduling

```
a();  
#pragma omp parallel for  
for (int i = 0; i < 16; ++i) {  
    w(i);  
}  
z();
```



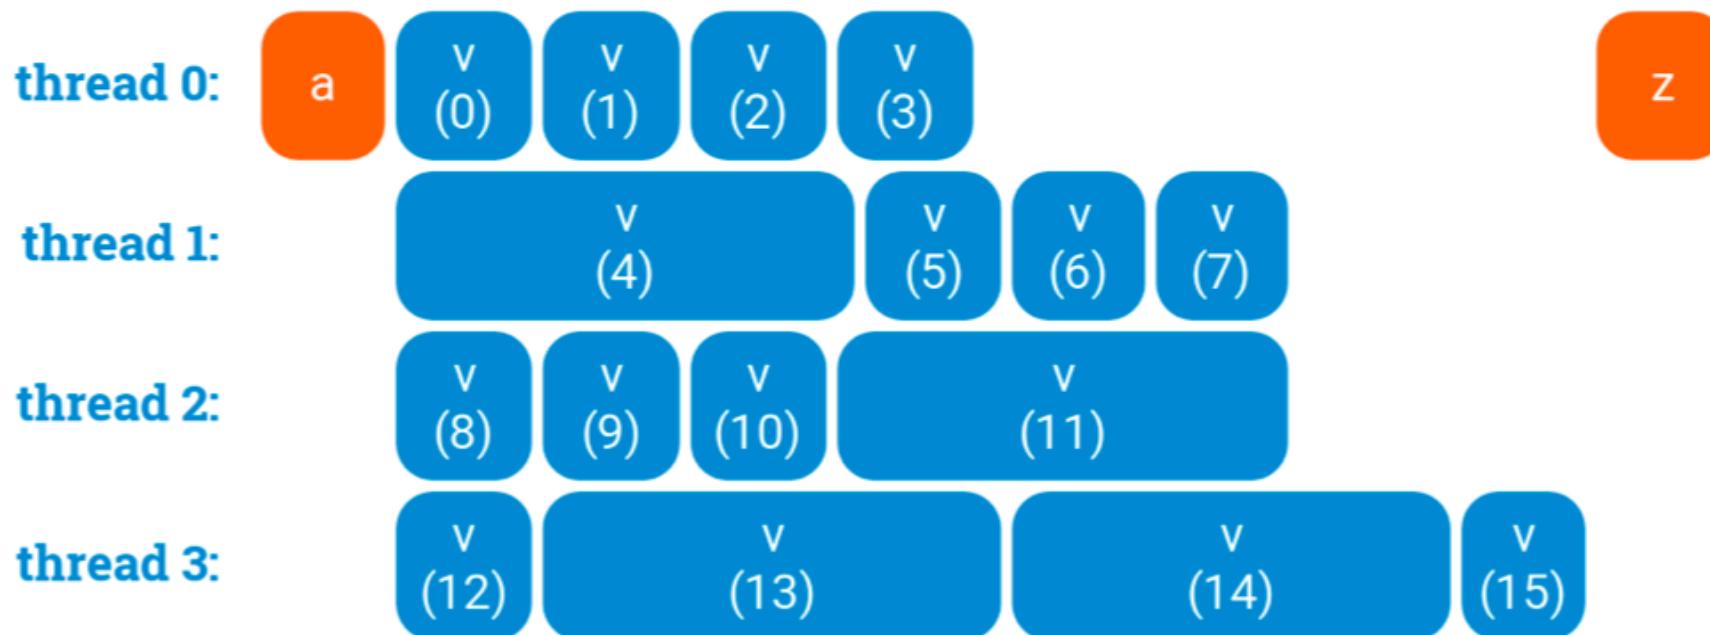
# for loop scheduling (static)

```
a();  
#pragma omp parallel for schedule(static,1)  
for (int i = 0; i < 16; ++i) {  
    w(i);  
}  
z();
```



# for loop scheduling

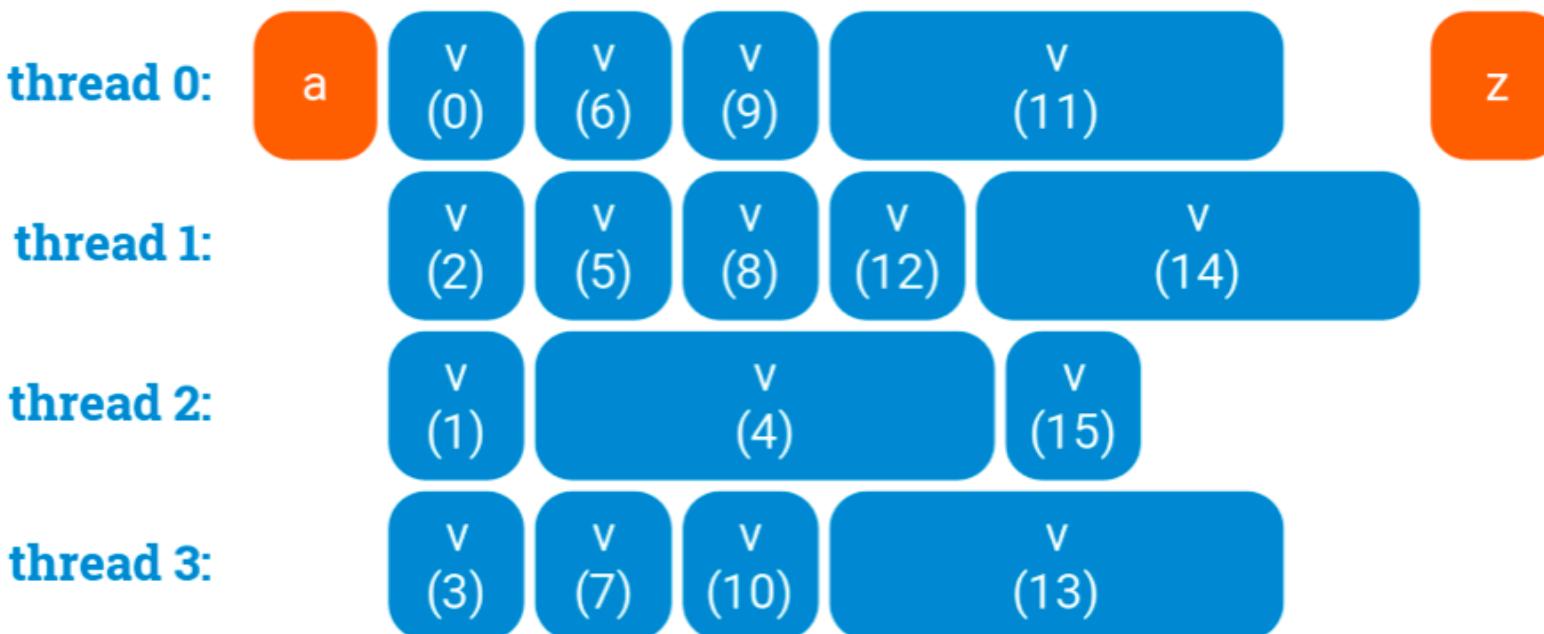
```
a();  
#pragma omp parallel for  
for (int i = 0; i < 16; ++i) {  
    v(i);  
}  
z();
```



# for loop scheduling (dynamic)

```
a();  
#pragma omp parallel for schedule(dynamic,1)  
for (int i = 0; i < 16; ++i) {  
    v(i);  
}  
z();
```

- dynamic scheduling is expensive
- no guarantee to be optimal

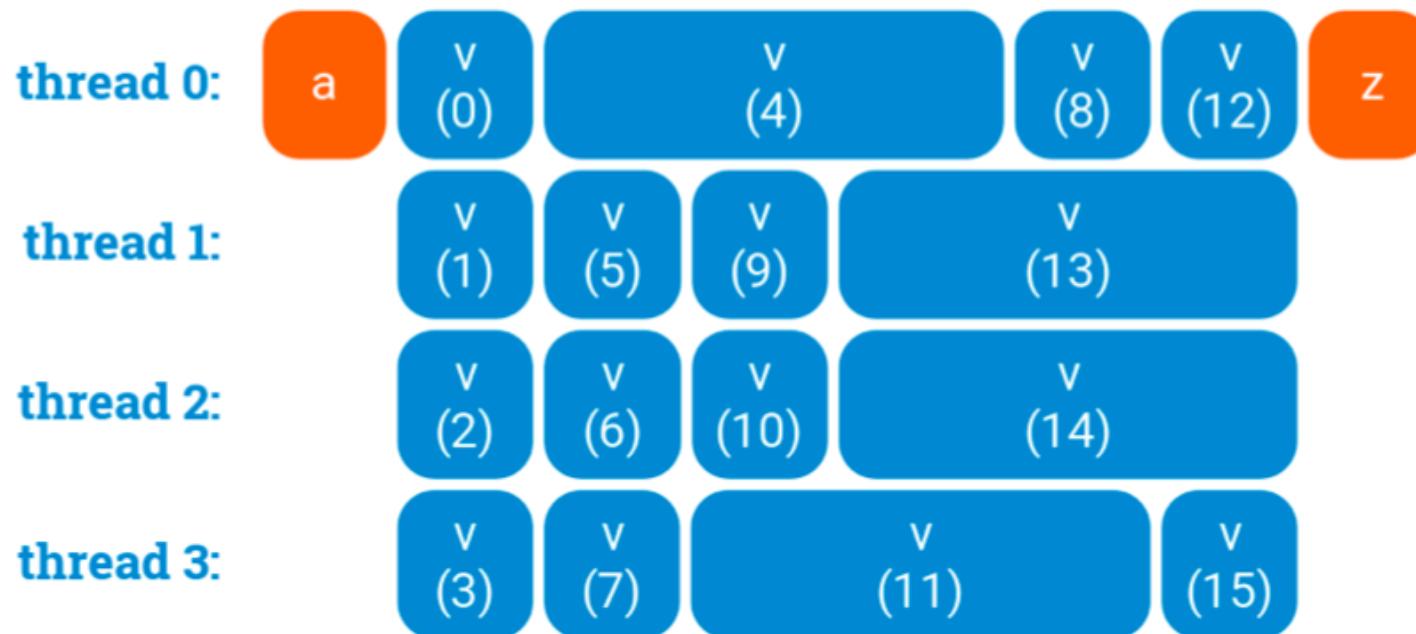


# for loop scheduling (dynamic)

```
a();  
#pragma omp parallel for schedule(dynamic,1)  
for (int i = 0; i < 16; ++i) {  
    v(i);  
}  
z();
```

- dynamic scheduling is expensive
- no guarantee to be optimal

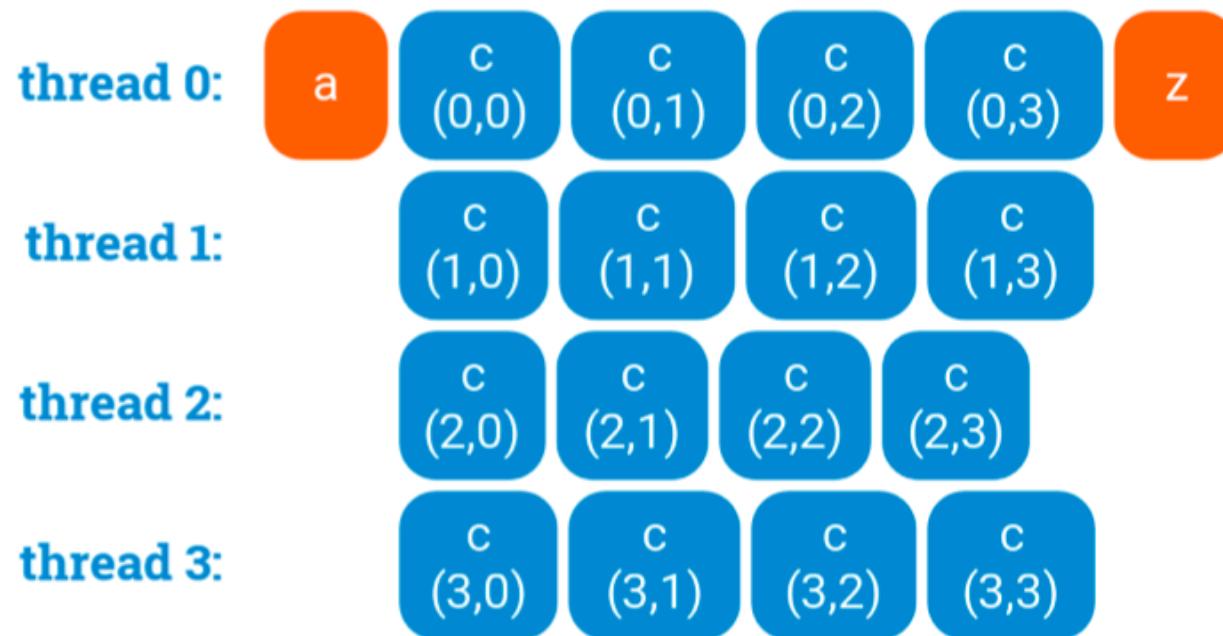
For example:



# Parallelizing nested loops

Often easy enough to parallelize the inner loop

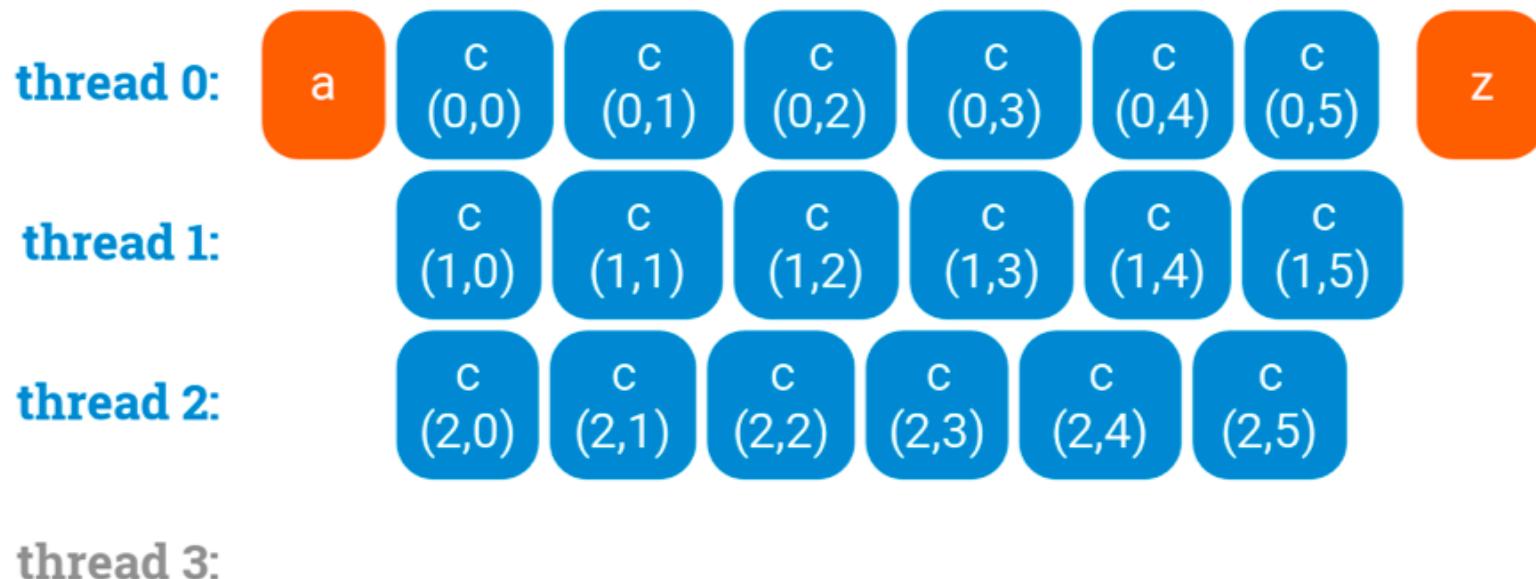
```
a();  
#pragma omp parallel for  
for (int i = 0; i < 4; ++i) {  
    for (int j = 0; j < 4; ++j) {  
        c(i, j);  
    }  
}  
z();
```



# Parallelizing nested loops

For example, the outer loop can be too short:

```
a();  
#pragma omp parallel for  
for (int i = 0; i < 3; ++i) {  
    for (int j = 0; j < 6; ++j) {  
        c(i, j);  
    }  
}  
z();
```

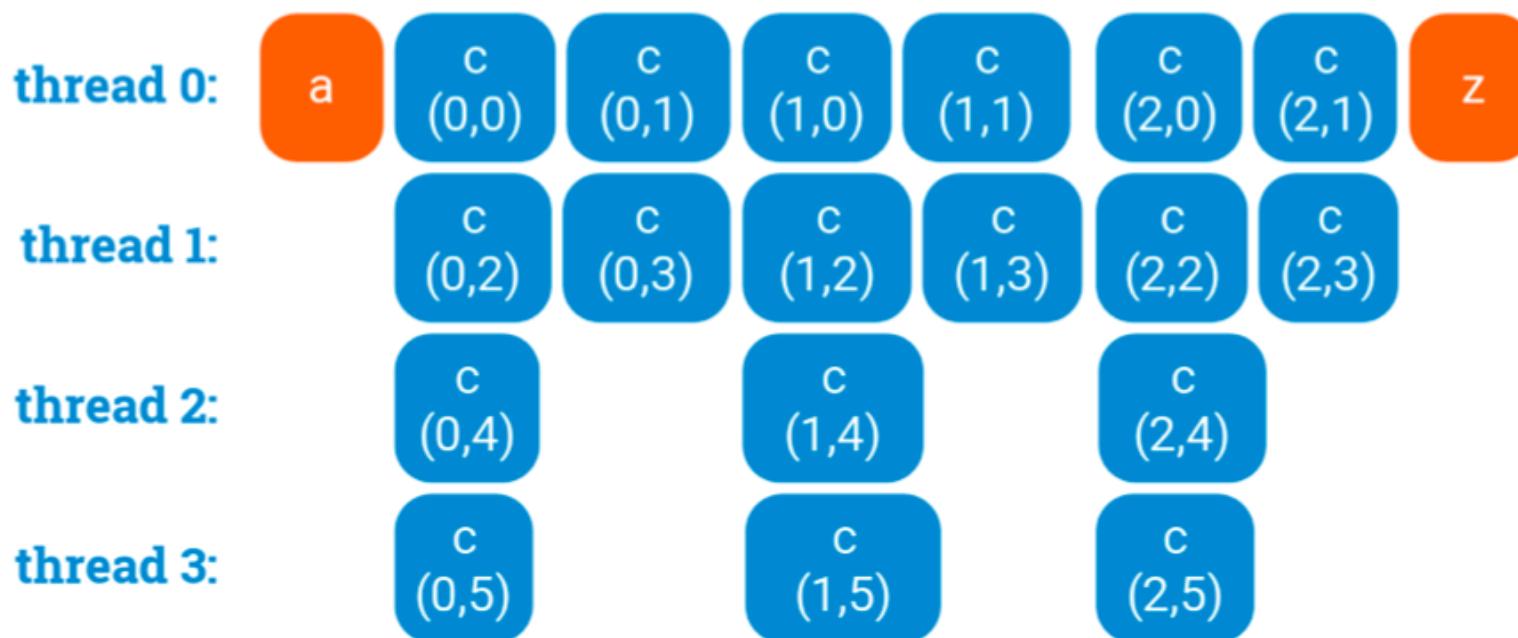


# Parallelizing nested loops (bad way)

Idea: parallelize the inner loop

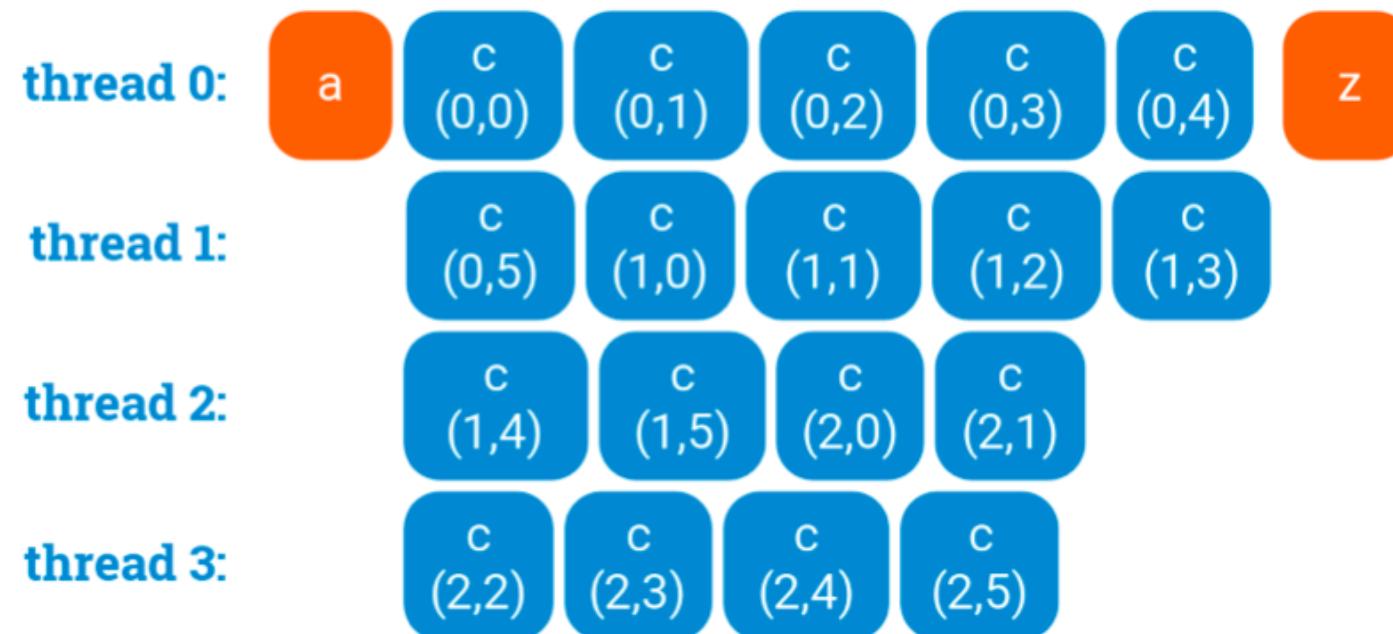
```
a();  
for (int i = 0; i < 3; ++i) {  
    #pragma omp parallel for  
    for (int j = 0; j < 6; ++j) {  
        c(i, j);  
    }  
}  
z();
```

However, no guarantee,  
that it will be any faster



# Parallelizing nested loops (good way 1)

```
a();  
#pragma omp parallel for  
for (int ij = 0; ij < 3 * 6; ++ij) {  
    c(ij / 6, ij % 6);  
}  
z();
```



# Parallelizing nested loops (good way 2)

```
a();  
#pragma omp parallel for collapse(2)  
for (int i = 0; i < 3; ++i) {  
    for (int j = 0; j < 6; ++j) {  
        c(i, j);  
    }  
}  
z();
```



# Race condition example

```
#include <stdio.h>
#include <omp.h>

int main(){
    int x = 2;

    #pragma omp parallel shared(x) num_threads(2)
    {
        if (1 == omp_get_thread_num()){
            x = 5;
        }
        else {
            printf("1: Thread %d has x = %d\n",omp_get_thread_num(),x);
        }

        #pragma omp barrier

        if (0 == omp_get_thread_num()) {
            printf("2: thread %d has x = %d\n",omp_get_thread_num(),x);
        }
        else {
            printf("3: thread %d has x = %x\n",omp_get_thread_num(),x);
        }

    }
    return 0;
}
```

# Synchronization in OpenMP

- Implicit **barrier** synchronization at end of parallel region (no explicit support for synch. subset of threads). Can invoke explicitly with `#pragma omp barrier`. All threads must see same sequence of work-sharing and barrier regions .
- **critical sections**: only one thread at a time in critical region with the same name. `#pragma omp critical [ (name) ]`
- **atomic** operation: protects updates to individual memory loc. Only simple expressions allowed. `#pragma omp atomic`
- **locks**: low-level run-time library routines (like mutex vars., semaphores)
- **flush** operation - forces the executing thread to make its values of shared data consistent with shared memory
- **master** (like single but not implied barrier at end), *ordered*, ...

At all these (implicit or explicit ) synchronization points OpenMP ensures that threads have consistent values of shared data.

# Atomic example

```
int sum = 0;  
#pragma omp parallel for shared(n,a,sum)  
{  
    for (i=0; i<n; i++) {  
        #pragma omp atomic  
        sum = sum + a[i];  
    }  
}
```

Better to use a *reduction* clause:

```
int sum = 0;  
#pragma omp parallel for shared(n,a)  \  
    reduction(+:sum)  
{  
    for (i=0; i<n; i++) {  
        sum += a[i];  
    }  
}
```

# Locks

Locks control access to shared resources. Up to implementation to use spin locks (busy waiting) or not.

- Lock variables must be accessed only through locking routines:

`omp_init_lock`   `omp_destroy_lock`  
`omp_set_lock`   `omp_unset_lock`   `omp_test_lock`

- In C, lock is a type `omp_lock_t` or `omp_nest_lock_t`  
(In Fortran lock variable is integer)
- initial state of lock is unlocked.
- `omp_set_lock (omp_lock_t *lock)` forces calling thread to wait until the specified lock is available.  
(Non-blocking version is `omp_test_lock`)

# Locks

```
#include <stdio.h>
#include <omp.h>

omp_lock_t my_lock;

int main() {
    omp_init_lock(&my_lock);

    #pragma omp parallel
    {
        int tid = omp_get_thread_num( );
        int i, j;

        for (i = 0; i < 5; ++i) {
            omp_set_lock(&my_lock);
            printf("Thread %d - starting locked region\n", tid);

            printf("Thread %d - ending locked region\n", tid);
            omp_unset_lock(&my_lock);
        }
    }

    omp_destroy_lock(&my_lock);
}
```

# Больше примеров использования:

- <https://www.openmp.org/resources/tutorials-articles/>



# Домашнее задание